



The Rust Programming Language

DARTH-REVAN

Originally created and published by Mozilla Foundation

https://github.com/Darth-Revan/rust-lang_Doc-LaTeX

March 31, 2016

Contents

1	Introduction	4
2	Getting Started	5
2.1	Installing Rust	5
2.2	Hello, World!	8
2.3	Hello, Cargo!	12
2.4	Closing Thoughts	16
3	Tutorial: Guessing Game	17
3.1	Set up	17
3.2	Processing a Guess	18
3.3	Generating a secret number	22
3.4	Compartmenting guesses	25
3.5	Looping	29
3.6	Complete	34
4	Syntax and Semantics	35
4.1	Variable Bindings	35
4.2	Functions	39
4.3	Primitive Types	44
4.4	Comments	49
4.5	If	50
4.6	Loops	52
4.7	Ownership	55
4.8	References and Borrowing	60
4.9	Lifetimes	67
4.10	Mutability	73
4.11	Structs	76
4.12	Enums	80
4.13	Match	82
4.14	Patterns	84
5	Effective Rust	90
6	Nightly Rust	91
7	Glossary	92
8	Syntax Index	94
9	Bibliography	95
9.1	Type system	95
9.2	Concurrency	95

9.3	Others	96
9.4	Papers about Rust	96

1 Introduction

Welcome! This book will teach you about the [Rust Programming Language](#). Rust is a systems programming language focused on three goals: safety, speed, and concurrency. It maintains these goals without having a garbage collector, making it a useful language for a number of use cases other languages aren't good at: embedding in other languages, programs with specific space and time requirements, and writing low-level code, like device drivers and operating systems. It improves on current languages targeting this space by having a number of compile-time safety checks that produce no runtime overhead, while eliminating all data races. Rust also aims to achieve 'zero-cost abstractions' even though some of these abstractions feel like those of a high-level language. Even then, Rust still allows precise control like a low-level language would.

"The Rust Programming Language" is split into chapters. This introduction is the first. After this:

- [Getting Started](#) - Set up your computer for Rust development.
- [Tutorial: Guessing Game](#) - Learn some Rust with a small project.
- [Syntax and Semantics](#) - Each bit of Rust, broken down into small chunks.
- [Effective Rust](#) - Higher-level concepts for writing excellent Rust code.
- [Nightly Rust](#) - Cutting-edge features that aren't in stable builds yet.
- [Glossary](#) - A reference of terms used in the book.
- [Bibliography](#) - Background on Rust's influences, papers about Rust.

Contributing

The source files from which this book is generated can be found on [GitHub](#).

2 Getting Started

This first chapter of the book will get us going with Rust and its tooling. First, we'll install Rust. Then, the classic 'Hello World' program. Finally, we'll talk about Cargo, Rust's build system and package manager.

2.1 Installing Rust

The first step to using Rust is to install it. Generally speaking, you'll need an Internet connection to run the commands in this section, as we'll be downloading Rust from the internet.

We'll be showing off a number of commands using a terminal, and those lines all start with `$`. We don't need to type in the `$s`, they are there to indicate the start of each command. We'll see many tutorials and examples around the web that follow this convention: `$` for commands run as our regular user, and `#` for commands we should be running as an administrator.

Platform support

The Rust compiler runs on, and compiles to, a great number of platforms, though not all platforms are equally supported. Rust's support levels are organized into three tiers, each with a different set of guarantees.

Platforms are identified by their "target triple" which is the string to inform the compiler what kind of output should be produced. The columns below indicate whether the corresponding component works on the specified platform.

Tier 1

Tier 1 platforms can be thought of as "guaranteed to build and work". Specifically they will each satisfy the following requirements:

- Automated testing is set up to run tests for the platform.
- Landing changes to the `rust-lang/rust` repository's master branch is gated on tests passing.
- Official release artifacts are provided for the platform.
- Documentation for how to use and how to build the platform is available.

Target	std	rustc	cargo	notes
<code>x86_64-pc-windows-msvc</code>	✓	✓	✓	64-bit MSVC (Windows 7+)
<code>i686-pc-windows-gnu</code>	✓	✓	✓	32-bit MinGW (Windows 7+)
<code>x86_64-pc-windows-gnu</code>	✓	✓	✓	64-bit MinGW (Windows 7+)
<code>i686-apple-darwin</code>	✓	✓	✓	32-bit OSX (10.7+, Lion+)
<code>x86_64-apple-darwin</code>	✓	✓	✓	64-bit OSX (10.7+, Lion+)
<code>i686-unknown-linux-gnu</code>	✓	✓	✓	32-bit Linux (2.6.18+)
<code>x86_64-unknown-linux-gnu</code>	✓	✓	✓	64-bit Linux (2.6.18+)

Tier 2

Tier 2 platforms can be thought of as “guaranteed to build”. Automated tests are not run so it’s not guaranteed to produce a working build, but platforms often work to quite a good degree and patches are always welcome! Specifically, these platforms are required to have each of the following:

- Automated building is set up, but may not be running tests.
- Landing changes to the `rust-lang/rust` repository’s master branch is gated on platforms **building**. Note that this means for some platforms only the standard library is compiled, but for others the full bootstrap is run.
- Official release artifacts are provided for the platform.

Target	std	rustc	cargo	notes
<code>i686-pc-windows-msvc</code>	✓	✓	✓	32-bit MSVC (Windows 7+)
<code>x86_64-unknown-linux-musl</code>	✓			64-bit Linux with MUSL
<code>arm-linux-androideabi</code>	✓			ARM Android
<code>arm-unknown-linux-gnueabi</code>	✓	✓		ARM Linux (2.6.18+)
<code>arm-unknown-linux-gnueabihf</code>	✓	✓		ARM Linux (2.6.18+)
<code>aarch64-unknown-linux-gnu</code>	✓			ARM64 Linux (2.6.18+)
<code>mips-unknown-linux-gnu</code>	✓			MIPS Linux (2.6.18+)
<code>mipsel-unknown-linux-gnu</code>	✓			MIPS (LE) Linux (2.6.18+)

Tier 3

Tier 3 platforms are those which Rust has support for, but landing changes is not gated on the platform either building or passing tests. Working builds for these platforms may be spotty as their reliability is often defined in terms of community contributions. Additionally, release artifacts and installers are not provided, but there may be community infrastructure producing these in unofficial locations.

Target	std	rustc	cargo	notes
i686-linux-android	✓			32-bit x86 Android
aarch64-linux-android	✓			ARM64 Android
powerpc-unknown-linux-gnu	✓			PowerPC Linux (2.6.18+)
i386-apple-ios	✓			32-bit x86 iOS
x86_64-apple-ios	✓			64-bit x86 iOS
armv7-apple-ios	✓			ARM iOS
armv7s-apple-ios	✓			ARM iOS
aarch64-apple-ios	✓			ARM64 iOS
i686-unknown-freebsd	✓	✓		32-bit FreeBSD
x86_64-unknown-freebsd	✓	✓		64-bit FreeBSD
x86_64-unknown-openbsd	✓	✓		64-bit OpenBSD
x86_64-unknown-netbsd	✓	✓		64-bit NetBSD
x86_64-unknown-bitrig	✓	✓		64-bit Bitrig
x86_64-unknown-dragonfly	✓	✓		64-bit DragonFlyBSD
x86_64-rumprun-netbsd	✓			64-bit NetBSD Rump Kernel
i686-pc-windows-msvc (XP)	✓			Windows XP support
x86_64-pc-windows-msvc (XP)	✓			Windows XP support

Note that this table can be expanded over time, this isn't the exhaustive set of tier 3 platforms that will ever be!

Installing on Linux or Mac

If we're on Linux or a Mac, all we need to do is open a terminal and type this:

```
$ curl -sSf https://static.rust-lang.org/rustup.sh | sh
```

This will download a script, and start the installation. If it all goes well, you'll see this appear:

```
Welcome to Rust.

This script will download the Rust compiler and its package manager,
Cargo, and install them to /usr/local. You may install elsewhere by
running this script with the --prefix=<path> option.

The installer will run under 'sudo' and may ask you for your password.
If you do not want the script to run 'sudo' then pass it the
--disable-sudo flag.

You may uninstall later by running /usr/local/lib/rustlib/uninstall.sh,
or by running this script again with the --uninstall flag.

Continue? (y/N)
```

From here, press **y** for 'yes', and then follow the rest of the prompts.

Installing on Windows

If you're on Windows, please download the appropriate [installer](#).

Uninstalling

Uninstalling Rust is as easy as installing it. On Linux or Mac, run the uninstall script:

```
$ sudo /usr/local/lib/rustlib/uninstall.sh
```

If we used the Windows installer, we can re-run the `.msi` and it will give us an uninstall option.

Troubleshooting

If we've got Rust installed, we can open up a shell, and type this:

```
$ rustc --version
```

You should see the version number, commit hash, and commit date.

If you do, Rust has been installed successfully! Congrats!

If you don't and you're on Windows, check that Rust is in your `%PATH%` system variable. If it isn't, run the installer again, select "Change" on the "Change, repair, or remove installation" page and ensure "Add to PATH" is installed on the local hard drive.

If not, there are a number of places where we can get help. The easiest is the [#rust IRC channel on irc.mozilla.org](#), which we can access through [Mibbit](#). Click that link, and we'll be chatting with other Rustaceans (a silly nickname we call ourselves) who can help us out. Other great resources include [the user's forum](#), and [Stack Overflow](#).

This installer also installs a copy of the documentation locally, so we can read it offline. On UNIX systems, `/usr/local/share/doc/rust` is the location. On Windows, it's in a `share/doc` directory, inside the directory to which Rust was installed.

2.2 Hello, World!

Now that you have Rust installed, we'll help you write your first Rust program. It's traditional when learning a new language to write a little program to print the text "Hello, world!" to the screen, and in this section, we'll follow that tradition.

The nice thing about starting with such a simple program is that you can quickly verify that your compiler is installed, and that it's working properly. Printing information to the screen is also a pretty common thing to do, so practicing it early on is good.

Note: This book assumes basic familiarity with the command line. Rust itself makes no specific demands about your editing, tooling, or where your code lives, so if you prefer an IDE to the command line, that's an option. You may want to check out SolidOak, which was built specifically with Rust in mind. There are a number of extensions in development by the community, and the Rust team ships plugins for various editors. Configuring your editor or IDE is out of the scope of this tutorial, so check the documentation for your specific setup.

Creating a Project File

First, make a file to put your Rust code in. Rust doesn't care where your code lives, but for this book, I suggest making a *projects* directory in your home directory, and keeping all your projects there. Open a terminal and enter the following commands to make a directory for this particular project:

```
$ mkdir ~/projects
$ cd ~/projects
$ mkdir hello_world
$ cd hello_world
```

Note: If you're on Windows and not using PowerShell, the `cd` may not work. Consult the documentation for your shell for more details.

Writing and Running a Rust Program

Next, make a new source file and call it *main.rs*. Rust files always end in a *.rs* extension. If you're using more than one word in your filename, use an underscore to separate them; for example, you'd use *hello_world.rs* rather than *helloworld.rs*.

Now open the *main.rs* file you just created, and type the following code:

```
fn main() {
    println!("Hello, world!");
}
```

Save the file, and go back to your terminal window. On Linux or OSX, enter the following commands:

```
$ rustc main.rs
$ ./main
Hello, world!
```

In Windows, replace `main` with `main.exe`. Regardless of your operating system, you should see the string `Hello, world!` print to the terminal. If you did, then congratulations! You've officially written a Rust program. That makes you a Rust programmer! Welcome.

Anatomy of a Rust Program

Now, let's go over what just happened in your "Hello, world!" program in detail. Here's the first piece of the puzzle:

```
fn main() {  
  
}
```

These lines define a *function* in Rust. The `main` function is special: it's the beginning of every Rust program. The first line says, "I'm declaring a function named `main` that takes no arguments and returns nothing." If there were arguments, they would go inside the parentheses (`(` and `)`), and because we aren't returning anything from this function, we can omit the return type entirely.

Also note that the function body is wrapped in curly braces (`{` and `}`). Rust requires these around all function bodies. It's considered good style to put the opening curly brace on the same line as the function declaration, with one space in between.

Inside the `main()` function:

```
println!("Hello, world!");
```

This line does all of the work in this little program: it prints text to the screen. There are a number of details that are important here. The first is that it's indented with four spaces, not tabs.

The second important part is the `println!()` line. This is calling a Rust *macro*, which is how metaprogramming is done in Rust. If it were calling a function instead, it would look like this: `println()` (without the `!`). We'll discuss Rust macros in more detail later, but for now you only need to know that when you see a `!` that means that you're calling a macro instead of a normal function.

Next is `"Hello, world!"` which is a *string*. Strings are a surprisingly complicated topic in a systems programming language, and this is a statically allocated string. We pass this string as an argument to `println!`, which prints the string to the screen. Easy enough!

The line ends with a semicolon (`;`). Rust is an [Expression-Oriented Language](#), which means that most things are expressions, rather than statements. The `;` indicates that this expression is over, and the next one is ready to begin. Most lines of Rust code end with a `;`.

Compiling and Running are Separate Steps

In “Writing and Running a Rust Program”, we showed you how to run a newly created program. We’ll break that process down and examine each step now.

Before running a Rust program, you have to compile it. You can use the Rust compiler by entering the `rustc` command and passing it the name of your source file, like this:

```
$ rustc main.rs
```

If you come from a C or C++ background, you’ll notice that this is similar to `gcc` or `clang`. After compiling successfully, Rust should output a binary executable, which you can see on Linux or OSX by entering the `ls` command in your shell as follows:

```
$ ls
main main.rs
```

On Windows, you’d enter:

```
$ dir
main.exe main.rs
```

This shows we have two files: the source code, with an `.rs` extension, and the executable (`main.exe` on Windows, `main` everywhere else). All that’s left to do from here is run the `main` or `main.exe` file, like this:

```
$ ./main # or main.exe on Windows
```

If `main.rs` were your “Hello, world!” program, this would print `Hello, world!` to your terminal.

If you come from a dynamic language like Ruby, Python, or JavaScript, you may not be used to compiling and running a program being separate steps. Rust is an *ahead-of-time compiled* language, which means that you can compile a program, give it to someone else, and they can run it even without Rust installed. If you give someone a `.rb` or `.py` or `.js` file, on the other hand, they need to have a Ruby, Python, or JavaScript implementation installed (respectively), but you only need one command to both compile and run your program. Everything is a tradeoff in language design.

Just compiling with `rustc` is fine for simple programs, but as your project grows, you’ll want to be able to manage all of the options your project has, and make it easy to share your code with other people and projects. Next, I’ll introduce you to a tool called Cargo, which will help you write real-world Rust programs.

2.3 Hello, Cargo!

Cargo is Rust's build system and package manager, and Rustaceans use Cargo to manage their Rust projects. Cargo manages three things: building your code, downloading the libraries your code depends on, and building those libraries. We call libraries your code needs 'dependencies' since your code depends on them.

The simplest Rust programs don't have any dependencies, so right now, you'd only use the first part of its functionality. As you write more complex Rust programs, you'll want to add dependencies, and if you start off using Cargo, that will be a lot easier to do.

As the vast, vast majority of Rust projects use Cargo, we will assume that you're using it for the rest of the book. Cargo comes installed with Rust itself, if you used the official installers. If you installed Rust through some other means, you can check if you have Cargo installed by typing:

```
$ cargo --version
```

into a terminal. If you see a version number, great! If you see an error like '**command not found**', then you should look at the documentation for the system in which you installed Rust, to determine if Cargo is separate.

Converting to Cargo

Let's convert the Hello World program to Cargo. To Cargo-fy a project, you need to do three things:

1. Put your source file in the right directory.
2. Get rid of the old executable (**main.exe** on Windows, **main** everywhere else) and make a new one.
3. Make a Cargo configuration file.

Let's get started!

Creating a new Executable and Source Directory

First, go back to your terminal, move to your *hello_world* directory, and enter the following commands:

```
$ mkdir src
$ mv main.rs src/main.rs
$ rm main # or 'del main.exe' on Windows
```

Cargo expects your source files to live inside a *src* directory, so do that first. This leaves the top-level project directory (in this case, *hello_world*) for READMEs, license information, and anything else not related to your code. In this way, using Cargo helps you keep your projects nice and tidy. There's a place for everything, and everything is in its place.

Now, copy *main.rs* to the *src* directory, and delete the compiled file you created with **rustc**. As usual, replace **main** with **main.exe** if you're on Windows.

This example retains **main.rs** as the source filename because it's creating an executable. If you wanted to make a library instead, you'd name the file **lib.rs**. This convention is used by Cargo to successfully compile your projects, but it can be overridden if you wish.

Creating a Configuration File

Next, create a new file inside your *hello_world* directory, and call it **Cargo.toml**.

Make sure to capitalize the **C** in *Cargo.toml*, or Cargo won't know what to do with the configuration file.

This file is in the **TOML** (Tom's Obvious, Minimal Language) format. TOML is similar to INI, but has some extra goodies, and is used as Cargo's configuration format.

Inside this file, type the following information:

```
[package]

name = "hello_world"
version = "0.0.1"
authors = [ "Your name <you@example.com>" ]
```

The first line, **[package]**, indicates that the following statements are configuring a package. As we add more information to this file, we'll add other sections, but for now, we only have the package configuration.

The other three lines set the three bits of configuration that Cargo needs to know to compile your program: its name, what version it is, and who wrote it.

Once you've added this information to the *Cargo.toml* file, save it to finish creating the configuration file.

Building and Running a Cargo Project

With your *Cargo.toml* file in place in your project's root directory, you should be ready to build and run your Hello World program! To do so, enter the following commands:

```
$ cargo build
   Compiling hello_world v0.0.1 (file:///home/yourname/projects/hello_world)
$ ./target/debug/hello_world
Hello, world!
```

Bam! If all goes well, **Hello, world!** should print to the terminal once more.

You just built a project with **cargo build** and ran it with **./target/debug/hello_world**, but you can actually do both in one step with **cargo run** as follows:

```
$ cargo run
   Running 'target/debug/hello_world'
Hello, world!
```

Notice that this example didn't re-build the project. Cargo figured out that the file hasn't changed, and so it just ran the binary. If you'd modified your source code, Cargo would have rebuilt the project before running it, and you would have seen something like this:

```
$ cargo run
   Compiling hello_world v0.0.1 (file:///home/yourname/projects/hello_world)
   Running 'target/debug/hello_world'
Hello, world!
```

Cargo checks to see if any of your project's files have been modified, and only rebuilds your project if they've changed since the last time you built it.

With simple projects, Cargo doesn't bring a whole lot over just using **rustc**, but it will become useful in future. This is especially true when you start using crates; these are synonymous with a 'library' or 'package' in other programming languages. For complex projects composed of multiple crates, it's much easier to let Cargo coordinate the build. Using Cargo, you can run **cargo build**, and it should work the right way.

Building for Release

When your project is finally ready for release, you can use **cargo build -release** to compile your project with optimizations. These optimizations make your Rust code run faster, but turning them on makes your program take longer to compile. This is why there are two different profiles, one for development, and one for building the final program you'll give to a user.

Running this command also causes Cargo to create a new file called *Cargo.lock*, which looks like this:

```
[root]
name = "hello_world"
version = "0.0.1"
```

Cargo uses the *Cargo.lock* file to keep track of dependencies in your application. This is the Hello World project's *Cargo.lock* file. This project doesn't have dependencies, so the file is a bit sparse. Realistically, you won't ever need to touch this file yourself; just let Cargo handle it.

That's it! If you've been following along, you should have successfully built `hello_world` with Cargo.

Even though the project is simple, it now uses much of the real tooling you'll use for the rest of your Rust career. In fact, you can expect to start virtually all Rust projects with some variation on the following commands:

```
$ git clone someurl.com/foo
$ cd foo
$ cargo build
```

Making a new Cargo Project the Easy Way

You don't have to go through that previous process every time you want to start a new project! Cargo can quickly make a bare-bones project directory that you can start developing in right away.

To start a new project with Cargo, enter `cargo new` at the command line:

```
$ cargo new hello_world --bin
```

This command passes `-bin` because the goal is to get straight to making an executable application, as opposed to a library. Executables are often called binaries (as in `/usr/bin`, if you're on a Unix system).

Cargo has generated two files and one directory for us: a `Cargo.toml` and a `src` directory with a `main.rs` file inside. These should look familiar, they're exactly what we created by hand, above.

This output is all you need to get started. First, open `Cargo.toml`. It should look something like this:

```
[package]

name = "hello_world"
version = "0.1.0"
```

```
authors = ["Your Name <you@example.com>"]
```

Cargo has populated *Cargo.toml* with reasonable defaults based on the arguments you gave it and your `git` global configuration. You may notice that Cargo has also initialized the `hello_world` directory as a `git` repository.

Here's what should be in `src/main.rs`:

```
fn main() {  
    println!("Hello, world!");  
}
```

Cargo has generated a “Hello World!” for you, and you’re ready to start coding!

Note: If you want to look at Cargo in more detail, check out the official [Cargo guide](#), which covers all of its features.

2.4 Closing Thoughts

This chapter covered the basics that will serve you well through the rest of this book, and the rest of your time with Rust. Now that you’ve got the tools down, we’ll cover more about the Rust language itself.

You have two options: Dive into a project with ‘Learn Rust’, or start from the bottom and work your way up with ‘Syntax and Semantics’. More experienced systems programmers will probably prefer ‘Learn Rust’, while those from dynamic backgrounds may enjoy either. Different people learn differently! Choose whatever’s right for you.

3 Tutorial: Guessing Game

Let's learn some Rust! For our first project, we'll implement a classic beginner programming problem: the guessing game. Here's how it works: Our program will generate a random integer between one and a hundred. It will then prompt us to enter a guess. Upon entering our guess, it will tell us if we're too low or too high. Once we guess correctly, it will congratulate us. Sounds good?

Along the way, we'll learn a little bit about Rust. The next chapter, 'Syntax and Semantics', will dive deeper into each part.

3.1 Set up

Let's set up a new project. Go to your projects directory. Remember how we had to create our directory structure and a `Cargo.toml` for `hello_world`? Cargo has a command that does that for us. Let's give it a shot:

```
$ cd ~/projects
$ cargo new guessing_game --bin
$ cd guessing_game
```

We pass the name of our project to `cargo new`, and then the `-bin` flag, since we're making a binary, rather than a library.

Check out the generated `Cargo.toml`:

```
[package]

name = "guessing_game"
version = "0.1.0"
authors = ["Your Name <you@example.com>"]
```

Cargo gets this information from your environment. If it's not correct, go ahead and fix that.

Finally, Cargo generated a 'Hello, world!' for us. Check out `src/main.rs`:

```
fn main() {
    println!("Hello, world!");
}
```

Let's try compiling what Cargo gave us:

```
$ cargo build
Compiling guessing_game v0.1.0 (file:///home/you/projects/guessing_game)
```

Excellent! Open up your `src/main.rs` again. We'll be writing all of our code in this file.

Before we move on, let me show you one more Cargo command: `run`. `cargo run` is kind of like `cargo build`, but it also then runs the produced executable. Try it out:

```
$ cargo run
   Compiling guessing_game v0.1.0 (file:///home/you/projects/guessing_game)
   Running 'target/debug/guessing_game'
Hello, world!
```

Great! The `run` command comes in handy when you need to rapidly iterate on a project. Our game is such a project, we need to quickly test each iteration before moving on to the next one.

3.2 Processing a Guess

Let's get to it! The first thing we need to do for our guessing game is allow our player to input a guess. Put this in your `src/main.rs`:

```
use std::io;

fn main() {
    println!("Guess the number!");

    println!("Please input your guess.");

    let mut guess = String::new();

    io::stdin().read_line(&mut guess)
        .expect("Failed to read line");

    println!("You guessed: {}", guess);
}
```

There's a lot here! Let's go over it, bit by bit.

```
use std::io;
```

We'll need to take user input, and then print the result as output. As such, we need the `io` library from the standard library. Rust only imports a few things by default into every program, the 'prelude'. If it's not in the prelude, you'll have to `use` it directly. There is also a second 'prelude', the `io prelude`, which serves a similar function: you import it, and it imports a number of useful, `io`-related things.

```
fn main() {
```

As you've seen before, the `main()` function is the entry point into your program. The `fn` syntax declares a new function, the `()`s indicate that there are no arguments, and `{` starts the body of the function. Because we didn't include a return type, it's assumed to be `()`, an empty tuple.

```
println!("Guess the number!");  
  
println!("Please input your guess.");
```

We previously learned that `println!()` is a macro that prints a string to the screen.

```
let mut guess = String::new();
```

Now we're getting interesting! There's a lot going on in this little line. The first thing to notice is that this is a `let` statement, which is used to create 'variable bindings'. They take this form:

```
let foo = bar;
```

This will create a new binding named `foo`, and bind it to the value `bar`. In many languages, this is called a 'variable', but Rust's variable bindings have a few tricks up their sleeves.

For example, they're immutable by default. That's why our example uses `mut`: it makes a binding mutable, rather than immutable. `let` doesn't take a name on the left hand side of the assignment, it actually accepts a 'pattern'. We'll use patterns later. It's easy enough to use for now:

```
let foo = 5; // immutable.  
let mut bar = 5; // mutable
```

Oh, and `//` will start a comment, until the end of the line. Rust ignores everything in comments.

So now we know that `let mut guess` will introduce a mutable binding named `guess`, but we have to look at the other side of the `=` for what it's bound to: `String::new()`.

`String` is a string type, provided by the standard library. A `String` is a growable, UTF-8 encoded bit of text.

The `::new()` syntax uses `::` because this is an 'associated function' of a particular type. That is to say, it's associated with `String` itself, rather than a particular instance of a `String`. Some languages call this a 'static method'.

This function is named `new()`, because it creates a new, empty `String`. You'll find a `new()` function on many types, as it's a common name for making a new value of some kind.

Let's move forward:

```
io::stdin().read_line(&mut guess)
    .expect("Failed to read line");
```

That's a lot more! Let's go bit-by-bit. The first line has two parts. Here's the first:

```
io::stdin()
```

Remember how we `used std::io` on the first line of the program? We're now calling an associated function on it. If we didn't `use std::io`, we could have written this line as `std::io::stdin()`.

This particular function returns a handle to the standard input for your terminal. More specifically, a `std::io::Stdin`.

The next part will use this handle to get input from the user:

```
.read_line(&mut guess)
```

Here, we call the `read_line()` method on our handle. Methods are like associated functions, but are only available on a particular instance of a type, rather than the type itself. We're also passing one argument to `read_line()`: `&mut guess`.

Remember how we bound `guess` above? We said it was mutable. However, `read_line` doesn't take a `String` as an argument: it takes a `&mut String`. Rust has a feature called 'references', which allows you to have multiple references to one piece of data, which can reduce copying. References are a complex feature, as one of Rust's major selling points is how safe and easy it is to use references. We don't need to know a lot of those details to finish our program right now, though. For now, all we need to know is that like `let` bindings, references are immutable by default. Hence, we need to write `&mut guess`, rather than `&guess`.

Why does `read_line()` take a mutable reference to a string? Its job is to take what the user types into standard input, and place that into a string. So it takes that string as an argument, and in order to add the input, it needs to be mutable.

But we're not quite done with this line of code, though. While it's a single line of text, it's only the first part of the single logical line of code:

```
.expect("Failed to read line");
```

When you call a method with the `.foo()` syntax, you may introduce a newline and other whitespace. This helps you split up long lines. We *could* have done:

```
io::stdin().read_line(&mut guess).expect("failed to read line");
```

But that gets hard to read. So we've split it up, three lines for three method calls. We already talked about `read_line()`, but what about `expect()`? Well, we already mentioned that `read_line()` puts what the user types into the `&mut String` we pass it. But it also returns a value: in this case, an `io::Result`. Rust has a number of types named `Result` in its standard library: a generic `Result`, and then specific versions for sub-libraries, like `io::Result`.

The purpose of these `Result` types is to encode error handling information. Values of the `Result` type, like any type, have methods defined on them. In this case, `io::Result` has an `expect()` method that takes a value it's called on, and if it isn't a successful one, panic!s with a message you passed it. A `panic!` like this will cause our program to crash, displaying the message.

If we leave off calling these two methods, our program will compile, but we'll get a warning:

```
$ cargo build
   Compiling guessing_game v0.1.0 (file:///home/you/projects/guessing_game)
src/main.rs:10:5: 10:39 warning: unused result which must be used,
#[warn(unused_must_use)] on by default
src/main.rs:10      io::stdin().read_line(&mut guess);
                    ~~~~~
```

Rust warns us that we haven't used the `Result` value. This warning comes from a special annotation that `io::Result` has. Rust is trying to tell you that you haven't handled a possible error. The right way to suppress the error is to actually write error handling. Luckily, if we want to crash if there's a problem, we can use these two little methods. If we can recover from the error somehow, we'd do something else, but we'll save that for a future project.

There's only one line of this first example left:

```
    println!("You guessed: {}", guess);
}
```

This prints out the string we saved our input in. The `{}`s are a placeholder, and so we pass it `guess` as an argument. If we had multiple `{}`s, we would pass multiple arguments:

```
let x = 5;
let y = 10;

println!("x and y: {} and {}", x, y);
```

Easy.

Anyway, that's the tour. We can run what we have with `cargo run`:

```
$ cargo run
  Compiling guessing_game v0.1.0 (file:///home/you/projects/guessing_game)
    Running 'target/debug/guessing_game'
Guess the number!
Please input your guess.
6
You guessed: 6
```

All right! Our first part is done: we can get input from the keyboard, and then print it back out.

3.3 Generating a secret number

Next, we need to generate a secret number. Rust does not yet include random number functionality in its standard library. The Rust team does, however, provide a [rand crate](#). A ‘crate’ is a package of Rust code. We’ve been building a ‘binary crate’, which is an executable. `rand` is a ‘library crate’, which contains code that’s intended to be used with other programs.

Using external crates is where Cargo really shines. Before we can write the code using `rand`, we need to modify our `Cargo.toml`. Open it up, and add these few lines at the bottom:

```
[dependencies]

rand="0.3.0"
```

The `[dependencies]` section of `Cargo.toml` is like the `[package]` section: everything that follows it is part of it, until the next section starts. Cargo uses the dependencies section to know what dependencies on external crates you have, and what versions you require. In this case, we’ve specified version `0.3.0`, which Cargo understands to be any release that’s compatible with this specific version. Cargo understands [Semantic Versioning](#), which is a standard for writing version numbers. A bare number like above is actually shorthand for `^0.3.0`, meaning “anything compatible with 0.3.0”. If we wanted to use only `0.3.0` exactly, we could say `rand="=0.3.0"` (note the two equal signs). And if we wanted to use the latest version we could use `*`. We could also use a range of versions. [Cargo’s documentation](#) contains more details.

Now, without changing any of our code, let’s build our project:

```
$ cargo build
  Updating registry 'https://github.com/rust-lang/crates.io-index'
  Downloading rand v0.3.8
  Downloading libc v0.1.6
    Compiling libc v0.1.6
    Compiling rand v0.3.8
```

```
Compiling guessing_game v0.1.0 (file:///home/you/projects/guessing_game)
```

(You may see different versions, of course.)

Lots of new output! Now that we have an external dependency, Cargo fetches the latest versions of everything from the registry, which is a copy of data from [Crates.io](https://crates.io). Crates.io is where people in the Rust ecosystem post their open source Rust projects for others to use.

After updating the registry, Cargo checks our `[dependencies]` and downloads any we don't have yet. In this case, while we only said we wanted to depend on `rand`, we've also grabbed a copy of `libc`. This is because `rand` depends on `libc` to work. After downloading them, it compiles them, and then compiles our project.

If we run `cargo build` again, we'll get different output:

```
$ cargo build
```

That's right, no output! Cargo knows that our project has been built, and that all of its dependencies are built, and so there's no reason to do all that stuff. With nothing to do, it simply exits. If we open up `src/main.rs` again, make a trivial change, and then save it again, we'll only see one line:

```
$ cargo build
Compiling guessing_game v0.1.0 (file:///home/you/projects/guessing_game)
```

So, we told Cargo we wanted any `0.3.x` version of `rand`, and so it fetched the latest version at the time this was written, `v0.3.8`. But what happens when next week, version `v0.3.9` comes out, with an important bugfix? While getting bugfixes is important, what if `0.3.9` contains a regression that breaks our code?

The answer to this problem is the `Cargo.lock` file you'll now find in your project directory. When you build your project for the first time, Cargo figures out all of the versions that fit your criteria, and then writes them to the `Cargo.lock` file. When you build your project in the future, Cargo will see that the `Cargo.lock` file exists, and then use that specific version rather than do all the work of figuring out versions again. This lets you have a repeatable build automatically. In other words, we'll stay at `0.3.8` until we explicitly upgrade, and so will anyone who we share our code with, thanks to the lock file.

What about when we do want to use `v0.3.9`? Cargo has another command, `update`, which says 'ignore the lock, figure out all the latest versions that fit what we've specified. If that works, write those versions out to the lock file'. But, by default, Cargo will only look for versions larger than `0.3.0` and smaller than `0.4.0`. If we want to move to

0.4.x, we'd have to update the `Cargo.toml` directly. When we do, the next time we `cargo build`, Cargo will update the index and re-evaluate our `rand` requirements.

There's a lot more to say about `Cargo` and its [ecosystem](#), but for now, that's all we need to know. Cargo makes it really easy to re-use libraries, and so Rustaceans tend to write smaller projects which are assembled out of a number of sub-packages.

Let's get on to actually *using* `rand`. Here's our next step:

```
extern crate rand;

use std::io;
use rand::Rng;

fn main() {
    println!("Guess the number!");

    let secret_number = rand::thread_rng().gen_range(1, 101);

    println!("The secret number is: {}", secret_number);

    println!("Please input your guess.");

    let mut guess = String::new();

    io::stdin().read_line(&mut guess)
        .expect("failed to read line");

    println!("You guessed: {}", guess);
}
```

The first thing we've done is change the first line. It now says `extern crate rand`. Because we declared `rand` in our [\[dependencies\]](#), we can use `extern crate` to let Rust know we'll be making use of it. This also does the equivalent of a `use rand`; as well, so we can make use of anything in the `rand` crate by prefixing it with `rand::`.

Next, we added another `use` line: `use rand::Rng`. We're going to use a method in a moment, and it requires that `Rng` be in scope to work. The basic idea is this: methods are defined on something called 'traits', and for the method to work, it needs the trait to be in scope. For more about the details, read the traits section.

There are two other lines we added, in the middle:

```
let secret_number = rand::thread_rng().gen_range(1, 101);

println!("The secret number is: {}", secret_number);
```

We use the `rand::thread_rng()` function to get a copy of the random number generator, which is local to the particular thread of execution we're in. Because we `use`

`rand::Rng`'d above, it has a `gen_range()` method available. This method takes two arguments, and generates a number between them. It's inclusive on the lower bound, but exclusive on the upper bound, so we need 1 and 101 to get a number ranging from one to a hundred.

The second line prints out the secret number. This is useful while we're developing our program, so we can easily test it out. But we'll be deleting it for the final version. It's not much of a game if it prints out the answer when you start it up!

Try running our new program a few times:

```
$ cargo run
  Compiling guessing_game v0.1.0 (file:///home/you/projects/guessing_game)
    Running 'target/debug/guessing_game'
Guess the number!
The secret number is: 7
Please input your guess.
4
You guessed: 4
$ cargo run
    Running 'target/debug/guessing_game'
Guess the number!
The secret number is: 83
Please input your guess.
5
You guessed: 5
```

Great! Next up: comparing our guess to the secret number.

3.4 Comparing guesses

Now that we've got user input, let's compare our guess to the secret number. Here's our next step, though it doesn't quite compile yet:

```
extern crate rand;

use std::io;
use std::cmp::Ordering;
use rand::Rng;

fn main() {
    println!("Guess the number!");

    let secret_number = rand::thread_rng().gen_range(1, 101);

    println!("The secret number is: {}", secret_number);

    println!("Please input your guess.");

    let mut guess = String::new();

    io::stdin().read_line(&mut guess)
        .expect("failed to read line");

    println!("You guessed: {}", guess);

    match guess.cmp(&secret_number) {
        Ordering::Less    => println!("Too small!"),
        Ordering::Greater => println!("Too big!"),
        Ordering::Equal   => println!("You win!"),
    }
}
```

A few new bits here. The first is another `use`. We bring a type called `std::cmp::Ordering` into scope. Then, five new lines at the bottom that use it:

```
match guess.cmp(&secret_number) {
    Ordering::Less    => println!("Too small!"),
    Ordering::Greater => println!("Too big!"),
    Ordering::Equal   => println!("You win!"),
}
```

The `cmp()` method can be called on anything that can be compared, and it takes a reference to the thing you want to compare it to. It returns the `Ordering` type we `used` earlier. We use a `match` statement to determine exactly what kind of `Ordering` it is. `Ordering` is an enum, short for ‘enumeration’, which looks like this:

```
enum Foo {
    Bar,
    Baz,
}
```

With this definition, anything of type `Foo` can be either a `Foo::Bar` or a `Foo::Baz`. We use the `::` to indicate the namespace for a particular `enum` variant.

The `Ordering` enum has three possible variants: `Less`, `Equal`, and `Greater`. The `match` statement takes a value of a type, and lets you create an ‘arm’ for each possible value. Since we have three types of `Ordering`, we have three arms:

```
match guess.cmp(&secret_number) {
    Ordering::Less    => println!("Too small!"),
    Ordering::Greater => println!("Too big!"),
    Ordering::Equal   => println!("You win!"),
}
```

If it’s `Less`, we print `Too small!`, if it’s `Greater`, `Too big!`, and if `Equal`, `You win!`. `match` is really useful, and is used often in Rust.

I did mention that this won’t quite compile yet, though. Let’s try it:

```
$ cargo build
   Compiling guessing_game v0.1.0 (file:///home/you/projects/guessing_game)
src/main.rs:28:21: 28:35 error: mismatched types:
  expected '&collections::string::String',
    found '&_'
(expected struct 'collections::string::String',
  found integral variable) [E0308]
src/main.rs:28      match guess.cmp(&secret_number) {
                        ~~~~~
error: aborting due to previous error
Could not compile 'guessing_game'.
```

Whew! This is a big error. The core of it is that we have ‘mismatched types’. Rust has a strong, static type system. However, it also has type inference. When we wrote `let guess = String::new()`, Rust was able to infer that `guess` should be a `String`, and so it doesn’t make us write out the type. And with our `secret_number`, there are a number of types which can have a value between one and a hundred: `i32`, a thirty-two-bit number, or `u32`, an unsigned thirty-two-bit number, or `i64`, a sixty-four-bit number or others. So far, that hasn’t mattered, and so Rust defaults to an `i32`. However, here, Rust doesn’t know how to compare the `guess` and the `secret_number`. They need to be the same type. Ultimately, we want to convert the `String` we read as input into a real number type, for comparison. We can do that with three more lines. Here’s our new program:

```
extern crate rand;

use std::io;
use std::cmp::Ordering;
use rand::Rng;

fn main() {
    println!("Guess the number!");

    let secret_number = rand::thread_rng().gen_range(1, 101);

    println!("The secret number is: {}", secret_number);

    println!("Please input your guess.");

    let mut guess = String::new();

    io::stdin().read_line(&mut guess)
        .expect("failed to read line");

    let guess: u32 = guess.trim().parse()
        .expect("Please type a number!");

    println!("You guessed: {}", guess);

    match guess.cmp(&secret_number) {
        Ordering::Less    => println!("Too small!"),
        Ordering::Greater => println!("Too big!"),
        Ordering::Equal   => println!("You win!"),
    }
}
```

The new three lines:

```
let guess: u32 = guess.trim().parse()
    .expect("Please type a number!");
```

Wait a minute, I thought we already had a `guess`? We do, but Rust allows us to ‘shadow’ the previous `guess` with a new one. This is often used in this exact situation, where `guess` starts as a `String`, but we want to convert it to an `u32`. Shadowing lets us re-use the `guess` name, rather than forcing us to come up with two unique names like `guess_str` and `guess`, or something else.

We bind `guess` to an expression that looks like something we wrote earlier:

```
guess.trim().parse()
```

Here, `guess` refers to the old `guess`, the one that was a `String` with our input in it. The `trim()` method on `Strings` will eliminate any white space at the beginning and end of our string. This is important, as we had to press the ‘return’ key to satisfy

`read_line()`. This means that if we type `5` and hit return, `guess` looks like this: `5`
`n`. The `n` represents ‘newline’, the enter key. `trim()` gets rid of this, leaving our string with only the `5`. The `parse()` method on strings parses a string into some kind of number. Since it can parse a variety of numbers, we need to give Rust a hint as to the exact type of number we want. Hence, `let guess: u32`. The colon (`:`) after `guess` tells Rust we’re going to annotate its type. `u32` is an unsigned, thirty-two bit integer. Rust has a number of built-in number types, but we’ve chosen `u32`. It’s a good default choice for a small positive number.

Just like `read_line()`, our call to `parse()` could cause an error. What if our string contained `A%`? There’d be no way to convert that to a number. As such, we’ll do the same thing we did with `read_line()`: use the `expect()` method to crash if there’s an error.

Let’s try our program out!

```
$ cargo run
  Compiling guessing_game v0.1.0 (file:///home/you/projects/guessing_game)
    Running 'target/guessing_game'
Guess the number!
The secret number is: 58
Please input your guess.
  76
You guessed: 76
Too big!
```

Nice! You can see I even added spaces before my guess, and it still figured out that I guessed 76. Run the program a few times, and verify that guessing the number works, as well as guessing a number too small.

Now we’ve got most of the game working, but we can only make one guess. Let’s change that by adding loops!

3.5 Looping

The `loop` keyword gives us an infinite loop. Let’s add that in:

```

extern crate rand;

use std::io;
use std::cmp::Ordering;
use rand::Rng;

fn main() {
    println!("Guess the number!");

    let secret_number = rand::thread_rng().gen_range(1, 101);

    println!("The secret number is: {}", secret_number);

    loop {
        println!("Please input your guess.");

        let mut guess = String::new();

        io::stdin().read_line(&mut guess)
            .expect("failed to read line");

        let guess: u32 = guess.trim().parse()
            .expect("Please type a number!");

        println!("You guessed: {}", guess);

        match guess.cmp(&secret_number) {
            Ordering::Less    => println!("Too small!"),
            Ordering::Greater => println!("Too big!"),
            Ordering::Equal   => println!("You win!"),
        }
    }
}

```

And try it out. But wait, didn't we just add an infinite loop? Yup. Remember our discussion about `parse()`? If we give a non-number answer, we'll **panic!** and quit. Observe:

```

$ cargo run
   Compiling guessing_game v0.1.0 (file:///home/you/projects/guessing_game)
   Running 'target/guessing_game'
Guess the number!
The secret number is: 59
Please input your guess.
45
You guessed: 45
Too small!
Please input your guess.
60
You guessed: 60
Too big!

```

```

Please input your guess.
59
You guessed: 59
You win!
Please input your guess.
quit
thread '<main>' panicked at 'Please type a number!'

```

Ha! **quit** actually quits. As does any other non-number input. Well, this is suboptimal to say the least. First, let's actually quit when you win the game:

```

extern crate rand;

use std::io;
use std::cmp::Ordering;
use rand::Rng;

fn main() {
    println!("Guess the number!");

    let secret_number = rand::thread_rng().gen_range(1, 101);

    println!("The secret number is: {}", secret_number);

    loop {
        println!("Please input your guess.");

        let mut guess = String::new();

        io::stdin().read_line(&mut guess)
            .expect("failed to read line");

        let guess: u32 = guess.trim().parse()
            .expect("Please type a number!");

        println!("You guessed: {}", guess);

        match guess.cmp(&secret_number) {
            Ordering::Less    => println!("Too small!"),
            Ordering::Greater => println!("Too big!"),
            Ordering::Equal   => {
                println!("You win!");
                break;
            }
        }
    }
}

```

By adding the **break** line after the **You win!**, we'll exit the loop when we win. Exiting the loop also means exiting the program, since it's the last thing in **main()**. We have only one more tweak to make: when someone inputs a non-number, we don't want to quit, we want to ignore it. We can do that like this:

```

extern crate rand;

use std::io;
use std::cmp::Ordering;
use rand::Rng;

fn main() {
    println!("Guess the number!");

    let secret_number = rand::thread_rng().gen_range(1, 101);

    println!("The secret number is: {}", secret_number);

    loop {
        println!("Please input your guess.");

        let mut guess = String::new();

        io::stdin().read_line(&mut guess)
            .expect("failed to read line");

        let guess: u32 = match guess.trim().parse() {
            Ok(num) => num,
            Err(_) => continue,
        };

        println!("You guessed: {}", guess);

        match guess.cmp(&secret_number) {
            Ordering::Less    => println!("Too small!"),
            Ordering::Greater => println!("Too big!"),
            Ordering::Equal   => {
                println!("You win!");
                break;
            }
        }
    }
}

```

These are the lines that changed:

```

let guess: u32 = match guess.trim().parse() {
    Ok(num) => num,
    Err(_) => continue,
};

```

This is how you generally move from ‘crash on error’ to ‘actually handle the returned by `parse()` is an `enum` like `Ordering`, but in this case, each variant has some data associated with it: `Ok` is a success, and `Err` is a failure. Each contains more information: the successfully parsed integer, or an error type. In this case, we `match` on `Ok(num)`, which sets the inner value of the `Ok` to the name `num`, and then we return it on the

right-hand side. In the `Err` case, we don't care what kind of error it is, so we use `_` instead of a name. This ignores the error, and `continue` causes us to go to the next iteration of the `loop`.

Now we should be good! Let's try:

```
$ cargo run
  Compiling guessing_game v0.1.0 (file:///home/you/projects/guessing_game)
  Running 'target/guessing_game'
Guess the number!
The secret number is: 61
Please input your guess.
10
You guessed: 10
Too small!
Please input your guess.
99
You guessed: 99
Too big!
Please input your guess.
foo
Please input your guess.
61
You guessed: 61
You win!
```

Awesome! With one tiny last tweak, we have finished the guessing game. Can you think of what it is? That's right, we don't want to print out the secret number. It was good for testing, but it kind of ruins the game. Here's our final source:

```
extern crate rand;

use std::io;
use std::cmp::Ordering;
use rand::Rng;

fn main() {
    println!("Guess the number!");

    let secret_number = rand::thread_rng().gen_range(1, 101);

    loop {
        println!("Please input your guess.");

        let mut guess = String::new();

        io::stdin().read_line(&mut guess)
            .expect("failed to read line");

        let guess: u32 = match guess.trim().parse() {
            Ok(num) => num,
            Err(_) => continue,
        };

        println!("You guessed: {}", guess);

        match guess.cmp(&secret_number) {
            Ordering::Less    => println!("Too small!"),
            Ordering::Greater => println!("Too big!"),
            Ordering::Equal   => {
                println!("You win!");
                break;
            }
        }
    }
}
```

3.6 Complete

At this point, you have successfully built the Guessing Game! Congratulations!

This first project showed you a lot: **let**, **match**, methods, associated functions, using external crates, and more. Our next project will show off even more.

4 Syntax and Semantics

This chapter breaks Rust down into small chunks, one for each concept.

If you'd like to learn Rust from the bottom up, reading this in order is a great way to do that.

These sections also form a reference for each concept, so if you're reading another tutorial and find something confusing, you can find it explained somewhere in here.

4.1 Variable Bindings

Virtually every non-'Hello World' Rust program uses *variable bindings*. They bind some value to a name, so it can be used later. `let` is used to introduce a binding, like this:

```
fn main() {  
    let x = 5;  
}
```

Putting `fn main() {` in each example is a bit tedious, so we'll leave that out in the future. If you're following along, make sure to edit your `main()` function, rather than leaving it off. Otherwise, you'll get an error.

Patterns

In many languages, a variable binding would be called a *variable*, but Rust's variable bindings have a few tricks up their sleeves. For example the left-hand side of a `let` expression is a 'pattern', not a variable name. This means we can do things like:

```
let (x, y) = (1, 2);
```

After this expression is evaluated, `x` will be one, and `y` will be two. Patterns are really powerful, and have their own section in the book. We don't need those features for now, so we'll keep this in the back of our minds as we go forward.

Type annotations

Rust is a statically typed language, which means that we specify our types up front, and they're checked at compile time. So why does our first example compile? Well, Rust has this thing called 'type inference'. If it can figure out what the type of something is, Rust doesn't require you to actually type it out.

We can add the type if we want to, though. Types come after a colon (`:`):

```
let x: i32 = 5;
```

If I asked you to read this out loud to the rest of the class, you'd say “**x** is a binding with the type **i32** and the value **five**.”

In this case we chose to represent **x** as a 32-bit signed integer. Rust has many different primitive integer types. They begin with **i** for signed integers and **u** for unsigned integers. The possible integer sizes are 8, 16, 32, and 64 bits.

In future examples, we may annotate the type in a comment. The examples will look like this:

```
fn main() {  
    let x = 5; // x: i32  
}
```

Note the similarities between this annotation and the syntax you use with **let**. Including these kinds of comments is not idiomatic Rust, but we'll occasionally include them to help you understand what the types that Rust infers are.

Mutability

By default, bindings are *immutable*. This code will not compile:

```
let x = 5;  
x = 10;
```

It will give you this error:

```
error: re-assignment of immutable variable 'x'  
    x = 10;  
    ~~~~~
```

If you want a binding to be mutable, you can use **mut**:

```
let mut x = 5; // mut x: i32  
x = 10;
```

There is no single reason that bindings are immutable by default, but we can think about it through one of Rust's primary focuses: safety. If you forget to say **mut**, the compiler will catch it, and let you know that you have mutated something you may not have intended to mutate. If bindings were mutable by default, the compiler would not be able to tell you this. If you *did* intend mutation, then the solution is quite easy: add **mut**.

There are other good reasons to avoid mutable state when possible, but they're out of the scope of this guide. In general, you can often avoid explicit mutation, and so it is preferable in Rust. That said, sometimes, mutation is what you need, so it's not verboten.

Initializing bindings

Rust variable bindings have one more aspect that differs from other languages: bindings are required to be initialized with a value before you're allowed to use them.

Let's try it out. Change your `src/main.rs` file to look like this:

```
fn main() {  
    let x: i32;  
  
    println!("Hello world!");  
}
```

You can use `cargo build` on the command line to build it. You'll get a warning, but it will still print "Hello, world!":

```
Compiling hello_world v0.0.1 (file:///home/you/projects/hello_world)  
src/main.rs:2:9: 2:10 warning: unused variable: 'x', #[warn(unused_variable)]  
on by default  
src/main.rs:2      let x: i32;  
                   ^
```

Rust warns us that we never use the variable binding, but since we never use it, no harm, no foul. Things change if we try to actually use this `x`, however. Let's do that. Change your program to look like this:

```
fn main() {  
    let x: i32;  
  
    println!("The value of x is: {}", x);  
}
```

And try to build it. You'll get an error:

```
$ cargo build  
Compiling hello_world v0.0.1 (file:///home/you/projects/hello_world)  
src/main.rs:4:39: 4:40 error: use of possibly uninitialized variable: 'x'  
src/main.rs:4      println!("The value of x is: {}", x);  
                   ^  
  
note: in expansion of format_args!  
<std macros>:2:23: 2:77 note: expansion site  
<std macros>:1:1: 3:2 note: in expansion of println!  
src/main.rs:4:5: 4:42 note: expansion site  
error: aborting due to previous error  
Could not compile 'hello_world'.
```

Rust will not let us use a value that has not been initialized. Next, let's talk about this stuff we've added to `println!`.

If you include two curly braces (`{}`, some call them moustaches...) in your string to print, Rust will interpret this as a request to interpolate some sort of value. *String*

interpolation is a computer science term that means "stick in the middle of a string." We add a comma, and then `x`, to indicate that we want `x` to be the value we're interpolating. The comma is used to separate arguments we pass to functions and macros, if you're passing more than one.

When you use the curly braces, Rust will attempt to display the value in a meaningful way by checking out its type. If you want to specify the format in a more detailed manner, there are a [wide number of options available](#). For now, we'll stick to the default: integers aren't very complicated to print.

Scope and shadowing

Let's get back to bindings. Variable bindings have a scope - they are constrained to live in a block they were defined in. A block is a collection of statements enclosed by `{` and `}`. Function definitions are also blocks! In the following example we define two variable bindings, `x` and `y`, which live in different blocks. `x` can be accessed from inside the `fn main() {}` block, while `y` can be accessed only from inside the inner block:

```
fn main() {
    let x: i32 = 17;
    {
        let y: i32 = 3;
        println!("The value of x is {} and value of y is {}", x, y);
    }
    println!("The value of x is {} and value of y is {}", x, y); // This won't work
}
```

The first `println!` would print "The value of x is 17 and the value of y is 3", but this example cannot be compiled successfully, because the second `println!` cannot access the value of `y`, since it is not in scope anymore. Instead we get this error:

```
$ cargo build
   Compiling hello v0.1.0 (file:///home/you/projects/hello_world)
main.rs:7:62: 7:63 error: unresolved name 'y'. Did you mean 'x'? [E0425]
main.rs:7      println!("The value of x is {} and value of y is {}", x, y); // This won't work
                                     ^
note: in expansion of format_args!
<std macros>:2:25: 2:56 note: expansion site
<std macros>:1:1: 2:62 note: in expansion of println!
<std macros>:3:1: 3:54 note: expansion site
<std macros>:1:1: 3:58 note: in expansion of println!
main.rs:7:5: 7:65 note: expansion site
main.rs:7:62: 7:63 help: run 'rustc --explain E0425' to see a detailed explanation
error: aborting due to previous error
Could not compile 'hello'.

To learn more, run the command again with --verbose.
```

Additionally, variable bindings can be shadowed. This means that a later variable

binding with the same name as another binding, that's currently in scope, will override the previous binding.

```
let x: i32 = 8;
{
    println!("{}", x); // Prints "8"
    let x = 12;
    println!("{}", x); // Prints "12"
}
println!("{}", x); // Prints "8"
let x = 42;
println!("{}", x); // Prints "42"
```

Shadowing and mutable bindings may appear as two sides of the same coin, but they are two distinct concepts that can't always be used interchangeably. For one, shadowing enables us to rebind a name to a value of a different type. It is also possible to change the mutability of a binding.

```
let mut x: i32 = 1;
x = 7;
let x = x; // x is now immutable and is bound to 7

let y = 4;
let y = "I can also be bound to text!"; // y is now of a different type
```

4.2 Functions

Every Rust program has at least one function, the `main` function:

```
fn main() {
}
```

This is the simplest possible function declaration. As we mentioned before, `fn` says 'this is a function', followed by the name, some parentheses because this function takes no arguments, and then some curly braces to indicate the body. Here's a function named `foo`:

```
fn foo() {
}
```

So, what about taking arguments? Here's a function that prints a number:

```
fn print_number(x: i32) {
    println!("x is: {}", x);
}
```

Here's a complete program that uses `print_number`:

```
fn main() {
    print_number(5);
}

fn print_number(x: i32) {
    println!("x is: {}", x);
}
```

As you can see, function arguments work very similar to `let` declarations: you add a type to the argument name, after a colon.

Here's a complete program that adds two numbers together and prints them:

```
fn main() {
    print_sum(5, 6);
}

fn print_sum(x: i32, y: i32) {
    println!("sum is: {}", x + y);
}
```

You separate arguments with a comma, both when you call the function, as well as when you declare it.

Unlike `let`, you must declare the types of function arguments. This does not work:

```
fn print_sum(x, y) {
    println!("sum is: {}", x + y);
}
```

You get this error:

```
expected one of '!', ':', or '@', found ')'
fn print_number(x, y) {
```

This is a deliberate design decision. While full-program inference is possible, languages which have it, like Haskell, often suggest that documenting your types explicitly is a best-practice. We agree that forcing functions to declare types while allowing for inference inside of function bodies is a wonderful sweet spot between full inference and no inference.

What about returning a value? Here's a function that adds one to an integer:

```
fn add_one(x: i32) -> i32 {
    x + 1
}
```

Rust functions return exactly one value, and you declare the type after an 'arrow', which is a dash (-) followed by a greater-than sign (>). The last line of a function determines what it returns. You'll note the lack of a semicolon here. If we added it in:


```
fn add_one(x: i32) -> i32 {
    x + 1;
}
```

We would get an error:

```
error: not all control paths return a value
fn add_one(x: i32) -> i32 {
    x + 1;
}

help: consider removing this semicolon:
    x + 1;
    ^
```

This reveals two interesting things about Rust: it is an expression-based language, and semicolons are different from semicolons in other ‘curly brace and semicolon’-based languages. These two things are related.

Expressions vs. Statements

Rust is primarily an expression-based language. There are only two kinds of statements, and everything else is an expression.

So what’s the difference? Expressions return a value, and statements do not. That’s why we end up with ‘not all control paths return a value’ here: the statement `x + 1;` doesn’t return a value. There are two kinds of statements in Rust: ‘declaration statements’ and ‘expression statements’. Everything else is an expression. Let’s talk about declaration statements first.

In some languages, variable bindings can be written as expressions, not statements. Like Ruby:

```
x = y = 5
```

In Rust, however, using `let` to introduce a binding is *not* an expression. The following will produce a compile-time error:

```
let x = (let y = 5); // expected identifier, found keyword 'let'
```

The compiler is telling us here that it was expecting to see the beginning of an expression, and a `let` can only begin a statement, not an expression.

Note that assigning to an already-bound variable (e.g. `y = 5`) is still an expression, although its value is not particularly useful. Unlike other languages where an assignment evaluates to the assigned value (e.g. `5` in the previous example), in Rust the value of an assignment is an empty tuple `()` because the assigned value can have only one owner, and any other returned value would be too surprising:

```
let mut y = 5;

let x = (y = 6); // x has the value '()', not '6'
```

The second kind of statement in Rust is the *expression statement*. Its purpose is to turn any expression into a statement. In practical terms, Rust's grammar expects statements to follow other statements. This means that you use semicolons to separate expressions from each other. This means that Rust looks a lot like most other languages that require you to use semicolons at the end of every line, and you will see semicolons at the end of almost every line of Rust code you see.

What is this exception that makes us say "almost"? You saw it already, in this code:

```
fn add_one(x: i32) -> i32 {
    x + 1
}
```

Our function claims to return an `i32`, but with a semicolon, it would return `()` instead. Rust realizes this probably isn't what we want, and suggests removing the semicolon in the error we saw before.

Early returns

But what about early returns? Rust does have a keyword for that, `return`:

```
fn foo(x: i32) -> i32 {
    return x;

    // we never run this code!
    x + 1
}
```

Using a `return` as the last line of a function works, but is considered poor style:

```
fn foo(x: i32) -> i32 {
    return x + 1;
}
```

The previous definition without `return` may look a bit strange if you haven't worked in an expression-based language before, but it becomes intuitive over time.

Diverging functions

Rust has some special syntax for 'diverging functions', which are functions that do not return:

```
fn diverges() -> ! {
    panic!("This function never returns!");
}
```

`panic!` is a macro, similar to `println!()` that we've already seen. Unlike `println!()`, `panic!()` causes the current thread of execution to crash with the given message. Because this function will cause a crash, it will never return, and so it has the type `!`, which is read 'diverges'.

If you add a main function that calls `diverges()` and run it, you'll get some output that looks like this:

```
thread '<main>' panicked at 'This function never returns!', hello.rs:2
```

If you want more information, you can get a backtrace by setting the `RUST_BACKTRACE` environment variable:

```
$ RUST_BACKTRACE=1 ./diverges
thread '<main>' panicked at 'This function never returns!', hello.rs:2
stack backtrace:
 1:      0x7f402773a829 - sys::backtrace::write::h0942de78b6c02817K8r
 2:      0x7f402773d7fc - panicking::on_panic::h3f23f9d0b5f4c91bu9w
 3:      0x7f402773960e - rt::unwind::begin_unwind_inner::h2844b8c5e81e79558Bw
 4:      0x7f4027738893 - rt::unwind::begin_unwind::h4375279447423903650
 5:      0x7f4027738809 - diverges::h2266b4c4b850236beaa
 6:      0x7f40277389e5 - main::h19bb1149c2f00ecfBaa
 7:      0x7f402773f514 - rt::unwind::try::try_fn::h13186883479104382231
 8:      0x7f402773d1d8 - __rust_try
 9:      0x7f402773f201 - rt::lang_start::ha172a3ce74bb453aK5w
10:      0x7f4027738a19 - main
11:      0x7f402694ab44 - __libc_start_main
12:      0x7f40277386c8 - <unknown>
13:             0x0 - <unknown>
```

`RUST_BACKTRACE` also works with Cargo's `run` command:

```
$ RUST_BACKTRACE=1 cargo run
Running 'target/debug/diverges'
thread '<main>' panicked at 'This function never returns!', hello.rs:2
stack backtrace:
 1:      0x7f402773a829 - sys::backtrace::write::h0942de78b6c02817K8r
 2:      0x7f402773d7fc - panicking::on_panic::h3f23f9d0b5f4c91bu9w
 3:      0x7f402773960e - rt::unwind::begin_unwind_inner::h2844b8c5e81e79558Bw
 4:      0x7f4027738893 - rt::unwind::begin_unwind::h4375279447423903650
 5:      0x7f4027738809 - diverges::h2266b4c4b850236beaa
 6:      0x7f40277389e5 - main::h19bb1149c2f00ecfBaa
 7:      0x7f402773f514 - rt::unwind::try::try_fn::h13186883479104382231
 8:      0x7f402773d1d8 - __rust_try
 9:      0x7f402773f201 - rt::lang_start::ha172a3ce74bb453aK5w
10:      0x7f4027738a19 - main
11:      0x7f402694ab44 - __libc_start_main
12:      0x7f40277386c8 - <unknown>
13:             0x0 - <unknown>
```

A diverging function can be used as any type:

```
let x: i32 = diverges();
let x: String = diverges();
```

Function pointers

We can also create variable bindings which point to functions:

```
let f: fn(i32) -> i32;
```

`f` is a variable binding which points to a function that takes an `i32` as an argument and returns an `i32`. For example:

```
fn plus_one(i: i32) -> i32 {
    i + 1
}

// without type inference
let f: fn(i32) -> i32 = plus_one;

// with type inference
let f = plus_one;
```

We can then use `f` to call the function:

```
let six = f(5);
```

4.3 Primitive Types

The Rust language has a number of types that are considered ‘primitive’. This means that they’re built-in to the language. Rust is structured in such a way that the standard library also provides a number of useful types built on top of these ones, as well, but these are the most primitive.

Booleans

Rust has a built in boolean type, named `bool`. It has two values, `true` and `false`:

```
let x = true;

let y: bool = false;
```

A common use of booleans is in if conditionals.

You can find more documentation for `bools` [in the standard library documentation](#).

char

The `char` type represents a single Unicode scalar value. You can create `chars` with a single tick: `'`

```
let x = 'x';
```

Unlike some other languages, this means that Rust's `char` is not a single byte, but four.

You can find more documentation for `chars` in [the standard library documentation](#).

Numeric types

Rust has a variety of numeric types in a few categories: signed and unsigned, fixed and variable, floating-point and integer.

These types consist of two parts: the category, and the size. For example, `u16` is an unsigned type with sixteen bits of size. More bits lets you have bigger numbers.

If a number literal has nothing to cause its type to be inferred, it defaults:

```
let x = 42; // x has type i32
let y = 1.0; // y has type f64
```

Here's a list of the different numeric types, with links to their documentation in the standard library:

- [i8](#)
- [i16](#)
- [i32](#)
- [i64](#)
- [u8](#)
- [u16](#)
- [u32](#)
- [u64](#)
- [isize](#)
- [usize](#)

- `f32`
- `f64`

Let's go over them by category:

Signed and Unsigned

Integer types come in two varieties: signed and unsigned. To understand the difference, let's consider a number with four bits of size. A signed, four-bit number would let you store numbers from `-8` to `+7`. Signed numbers use “two's complement representation”. An unsigned four bit number, since it does not need to store negatives, can store values from `0` to `+15`.

Unsigned types use a `u` for their category, and signed types use `i`. The `i` is for ‘integer’. So `u8` is an eight-bit unsigned number, and `i8` is an eight-bit signed number.

Fixed size types

Fixed size types have a specific number of bits in their representation. Valid bit sizes are `8`, `16`, `32`, and `64`. So, `u32` is an unsigned, 32-bit integer, and `i64` is a signed, 64-bit integer.

Variable sized types

Rust also provides types whose size depends on the size of a pointer of the underlying machine. These types have ‘size’ as the category, and come in signed and unsigned varieties. This makes for two types: `isize` and `usize`.

Floating-point types

Rust also has two floating point types: `f32` and `f64`. These correspond to IEEE-754 single and double precision numbers.

Arrays

Like many programming languages, Rust has list types to represent a sequence of things. The most basic is the *array*, a fixed-size list of elements of the same type. By default, arrays are immutable.

```
let a = [1, 2, 3]; // a: [i32; 3]
let mut m = [1, 2, 3]; // m: [i32; 3]
```

Arrays have type `[T; N]`. We'll talk about this `T` notation in the generics section. The `N` is a compile-time constant, for the length of the array.

There's a shorthand for initializing each element of an array to the same value. In this example, each element of `a` will be initialized to `0`:

```
let a = [0; 20]; // a: [i32; 20]
```

You can get the number of elements in an array `a` with `a.len()`:

```
let a = [1, 2, 3];  
  
println!("a has {} elements", a.len());
```

You can access a particular element of an array with *subscript notation*:

```
let names = ["Graydon", "Brian", "Niko"]; // names: [&str; 3]  
  
println!("The second name is: {}", names[1]);
```

Subscripts start at zero, like in most programming languages, so the first name is `names[0]` and the second name is `names[1]`. The above example prints **The second name is: Brian**. If you try to use a subscript that is not in the array, you will get an error: array access is bounds-checked at run-time. Such errant access is the source of many bugs in other systems programming languages.

You can find more documentation for [arrays in the standard library documentation](#).

Slices

A ‘slice’ is a reference to (or “view” into) another data structure. They are useful for allowing safe, efficient access to a portion of an array without copying. For example, you might want to reference only one line of a file read into memory. By nature, a slice is not created directly, but from an existing variable binding. Slices have a defined length, can be mutable or immutable.

Slicing syntax

You can use a combo of `&` and `[]` to create a slice from various things. The `&` indicates that slices are similar to references, which we will cover in detail later in this section. The `[]`s, with a range, let you define the length of the slice:

```
let a = [0, 1, 2, 3, 4];  
let complete = &a[..]; // A slice containing all of the elements in a  
let middle = &a[1..4]; // A slice of a: only the elements 1, 2, and 3
```

Slices have type `&[T]`. We'll talk about that `T` when we cover generics.

You can find more documentation for slices [in the standard library documentation](#).

str

Rust's `str` type is the most primitive string type. As an unsized type, it's not very useful by itself, but becomes useful when placed behind a reference, like `&str`. We'll elaborate further when we cover Strings and references.

You can find more documentation for `str` in the [standard library documentation](#).

Tuples

A tuple is an ordered list of fixed size. Like this:

```
let x = (1, "hello");
```

The parentheses and commas form this two-length tuple. Here's the same code, but with the type annotated:

```
let x: (i32, &str) = (1, "hello");
```

As you can see, the type of a tuple looks like the tuple, but with each position having a type name rather than the value. Careful readers will also note that tuples are heterogeneous: we have an `i32` and a `&str` in this tuple. In systems programming languages, strings are a bit more complex than in other languages. For now, read `&str` as a *string slice*, and we'll learn more soon.

You can assign one tuple into another, if they have the same contained types and *Arity*. Tuples have the same arity when they have the same length.

```
let mut x = (1, 2); // x: (i32, i32)
let y = (2, 3); // y: (i32, i32)

x = y;
```

You can access the fields in a tuple through a *destructuring let*. Here's an example:

```
let (x, y, z) = (1, 2, 3);

println!("x is {}", x);
```

Remember before when I said the left-hand side of a `let` statement was more powerful than assigning a binding? Here we are. We can put a pattern on the left-hand side of the `let`, and if it matches up to the right-hand side, we can assign multiple bindings at once. In this case, `let` “destructures” or “breaks up” the tuple, and assigns the bits to three bindings.

This pattern is very powerful, and we'll see it repeated more later.

You can disambiguate a single-element tuple from a value in parentheses with a comma:

```
(0,); // single-element tuple
(0); // zero in parentheses
```

Tuple Indexing

You can also access fields of a tuple with indexing syntax:

```
let tuple = (1, 2, 3);

let x = tuple.0;
let y = tuple.1;
let z = tuple.2;

println!("x is {}", x);
```

Like array indexing, it starts at zero, but unlike array indexing, it uses a `.`, rather than `[]`s.

You can find more documentation for tuples [in the standard library documentation](#).

Functions

Functions also have a type! They look like this:

```
fn foo(x: i32) -> i32 { x }

let x: fn(i32) -> i32 = foo;
```

In this case, `x` is a ‘function pointer’ to a function that takes an `i32` and returns an `i32`.

4.4 Comments

Now that we have some functions, it’s a good idea to learn about comments. Comments are notes that you leave to other programmers to help explain things about your code. The compiler mostly ignores them.

Rust has two kinds of comments that you should care about: *line comments* and *doc comments*.

```
// Line comments are anything after '//' and extend to the end of the line.

let x = 5; // this is also a line comment.

// If you have a long explanation for something, you can put line comments next
// to each other. Put a space between the // and your comment so that it's
// more readable.
```

The other kind of comment is a doc comment. Doc comments use `///` instead of `//`, and support Markdown notation inside:

```
/// Adds one to the number given.
///
/// # Examples
///
/// ```
/// let five = 5;
///
/// assert_eq!(6, add_one(5));
/// # fn add_one(x: i32) -> i32 {
/// #     x + 1
/// # }
/// ```
fn add_one(x: i32) -> i32 {
    x + 1
}
```

There is another style of doc comment, `//!`, to comment containing items (e.g. crates, modules or functions), instead of the items following it. Commonly used inside crates root (lib.rs) or modules root (mod.rs):

```
//! # The Rust Standard Library
//!
//! The Rust Standard Library provides the essential runtime
//! functionality for building portable Rust software.
```

When writing doc comments, providing some examples of usage is very, very helpful. You'll notice we've used a new macro here: `assert_eq!`. This compares two values, and `panic!`s if they're not equal to each other. It's very helpful in documentation. There's another macro, `assert!`, which `panic!`s if the value passed to it is `false`.

You can use the `rustdoc` tool to generate HTML documentation from these doc comments, and also to run the code examples as tests!

4.5 If

Rust's take on `if` is not particularly complex, but it's much more like the `if` you'll find in a dynamically typed language than in a more traditional systems language. So let's talk about it, to make sure you grasp the nuances.

`if` is a specific form of a more general concept, the ‘branch’. The name comes from a branch in a tree: a decision point, where depending on a choice, multiple paths can be taken.

In the case of `if`, there is one choice that leads down two paths:

```
let x = 5;

if x == 5 {
    println!("x is five!");
}
```

If we changed the value of `x` to something else, this line would not print. More specifically, if the expression after the `if` evaluates to `true`, then the block is executed. If it's `false`, then it is not.

If you want something to happen in the `false` case, use an `else`:

```
let x = 5;

if x == 5 {
    println!("x is five!");
} else {
    println!("x is not five :(");
}
```

If there is more than one case, use an `else if`:

```
let x = 5;

if x == 5 {
    println!("x is five!");
} else if x == 6 {
    println!("x is six!");
} else {
    println!("x is not five or six :(");
}
```

This is all pretty standard. However, you can also do this:

```
let x = 5;

let y = if x == 5 {
    10
} else {
    15
}; // y: i32
```

Which we can (and probably should) write like this:

```
let x = 5;

let y = if x == 5 { 10 } else { 15 }; // y: i32
```

This works because `if` is an expression. The value of the expression is the value of the last expression in whichever branch was chosen. An `if` without an `else` always results in `()` as the value.

4.6 Loops

Rust currently provides three approaches to performing some kind of iterative activity. They are: `loop`, `while` and `for`. Each approach has its own set of uses.

loop

The infinite `loop` is the simplest form of loop available in Rust. Using the keyword `loop`, Rust provides a way to loop indefinitely until some terminating statement is reached. Rust's infinite `loops` look like this:

```
loop {
    println!("Loop forever!");
}
```

while

Rust also has a `while` loop. It looks like this:

```
let mut x = 5; // mut x: i32
let mut done = false; // mut done: bool

while !done {
    x += x - 3;

    println!("{}", x);

    if x % 5 == 0 {
        done = true;
    }
}
```

`while` loops are the correct choice when you're not sure how many times you need to loop.

If you need an infinite loop, you may be tempted to write this:

```
while true {
```

However, `loop` is far better suited to handle this case:

```
loop {
```

Rust’s control-flow analysis treats this construct differently than a `while true`, since we know that it will always loop. In general, the more information we can give to the compiler, the better it can do with safety and code generation, so you should always prefer `loop` when you plan to loop infinitely.

for

The `for` loop is used to loop a particular number of times. Rust’s `for` loops work a bit differently than in other systems languages, however. Rust’s `for` loop doesn’t look like this “C-style” `for` loop:

```
for (x = 0; x < 10; x++) {  
    printf( "%d\n", x );  
}
```

Instead, it looks like this:

```
for x in 0..10 {  
    println!("{}", x); // x: i32  
}
```

In slightly more abstract terms,

```
for var in expression {  
    code  
}
```

The expression is an item that can be converted into an iterator using `.iter`. The iterator gives back a series of elements. Each element is one iteration of the loop. That value is then bound to the name `var`, which is valid for the loop body. Once the body is over, the next value is fetched from the iterator, and we loop another time. When there are no more values, the `for` loop is over.

In our example, `0..10` is an expression that takes a start and an end position, and gives an iterator over those values. The upper bound is exclusive, though, so our loop will print `0` through `9`, not `10`.

Rust does not have the “C-style” `for` loop on purpose. Manually controlling each element of the loop is complicated and error prone, even for experienced C developers.

Enumerate

When you need to keep track of how many times you already looped, you can use the `.enumerate()` function.

On ranges:

```
for (i,j) in (5..10).enumerate() {  
    println!("i = {} and j = {}", i, j);  
}
```

Outputs:

```
i = 0 and j = 5  
i = 1 and j = 6  
i = 2 and j = 7  
i = 3 and j = 8  
i = 4 and j = 9
```

Don't forget to add the parentheses around the range.

On iterators:

```
for (linenumber, line) in lines.enumerate() {  
    println!("{: {}:", linenumber, line);  
}
```

Outputs:

```
0: Content of line one  
1: Content of line two  
2: Content of line three  
3: Content of line four
```

Ending iteration early

Let's take a look at that **while** loop we had earlier:

```
let mut x = 5;  
let mut done = false;  
  
while !done {  
    x += x - 3;  
  
    println!("{}", x);  
  
    if x % 5 == 0 {  
        done = true;  
    }  
}
```

We had to keep a dedicated **mut** boolean variable binding, **done**, to know when we should exit out of the loop. Rust has two keywords to help us with modifying iteration: **break** and **continue**.

In this case, we can write the loop in a better way with **break**:

```
let mut x = 5;

loop {
    x += x - 3;

    println!("{}", x);

    if x % 5 == 0 { break; }
}
```

We now loop forever with **loop** and use **break** to break out early. Issuing an explicit **return** statement will also serve to terminate the loop early.

continue is similar, but instead of ending the loop, goes to the next iteration. This will only print the odd numbers:

```
for x in 0..10 {
    if x % 2 == 0 { continue; }

    println!("{}", x);
}
```

Loop labels

You may also encounter situations where you have nested loops and need to specify which one your **break** or **continue** statement is for. Like most other languages, by default a **break** or **continue** will apply to innermost loop. In a situation where you would like to a **break** or **continue** for one of the outer loops, you can use labels to specify which loop the **break** or **continue** statement applies to. This will only print when both **x** and **y** are odd:

```
'outer: for x in 0..10 {
    'inner: for y in 0..10 {
        if x % 2 == 0 { continue 'outer; } // continues the loop over x
        if y % 2 == 0 { continue 'inner; } // continues the loop over y
        println!("x: {}, y: {}", x, y);
    }
}
```

4.7 Ownership

This guide is one of three presenting Rust's ownership system. This is one of Rust's most unique and compelling features, with which Rust developers should become quite acquainted. Ownership is how Rust achieves its largest goal, memory safety. There are a few distinct concepts, each with its own chapter:

- ownership, which you're reading now
- borrowing, and their associated feature 'references'
- lifetimes, an advanced concept of borrowing

These three chapters are related, and in order. You'll need all three to fully understand the ownership system.

Meta

Before we get to the details, two important notes about the ownership system.

Rust has a focus on safety and speed. It accomplishes these goals through many 'zero-cost abstractions', which means that in Rust, abstractions cost as little as possible in order to make them work. The ownership system is a prime example of a zero-cost abstraction. All of the analysis we'll talk about in this guide is *done at compile time*. You do not pay any run-time cost for any of these features.

However, this system does have a certain cost: learning curve. Many new users to Rust experience something we like to call 'fighting with the borrow checker', where the Rust compiler refuses to compile a program that the author thinks is valid. This often happens because the programmer's mental model of how ownership should work doesn't match the actual rules that Rust implements. You probably will experience similar things at first. There is good news, however: more experienced Rust developers report that once they work with the rules of the ownership system for a period of time, they fight the borrow checker less and less.

With that in mind, let's learn about ownership.

Ownership

[Variable Bindings](#) have a property in Rust: they 'have ownership' of what they're bound to. This means that when a binding goes out of scope, Rust will free the bound resources. For example:

```
fn foo() {  
    let v = vec![1, 2, 3];  
}
```

When **v** comes into scope, a new [vector] is created, and it allocates space on the heap for each of its elements. When **v** goes out of scope at the end of **foo()**, Rust will clean up everything related to the vector, even the heap-allocated memory. This happens deterministically, at the end of the scope.

We'll cover vectors in detail later in this chapter; we only use them here as an example of a type that allocates space on the heap at runtime. They behave like arrays, except their size may change by `push()`ing more elements onto them.

Vectors have a generic type `Vec<T>`, so in this example `v` will have type `Vec<i32>`. We'll cover generics in detail later in this chapter.

Move semantics

There's some more subtlety here, though: Rust ensures that there is *exactly one* binding to any given resource. For example, if we have a vector, we can assign it to another binding:

```
let v = vec![1, 2, 3];

let v2 = v;
```

But, if we try to use `v` afterwards, we get an error:

```
let v = vec![1, 2, 3];

let v2 = v;

println!("v[0] is: {}", v[0]);
```

It looks like this:

```
error: use of moved value: 'v'
println!("v[0] is: {}", v[0]);
                        ^
```

A similar thing happens if we define a function which takes ownership, and try to use something after we've passed it as an argument:

```
fn take(v: Vec<i32>) {
    // what happens here isn't important.
}

let v = vec![1, 2, 3];

take(v);

println!("v[0] is: {}", v[0]);
```

Same error: 'use of moved value'. When we transfer ownership to something else, we say that we've 'moved' the thing we refer to. You don't need some sort of special annotation here, it's the default thing that Rust does.

The details

The reason that we cannot use a binding after we've moved it is subtle, but important. When we write code like this:

```
let v = vec![1, 2, 3];  
  
let v2 = v;
```

The first line allocates memory for the vector object, `v`, and for the data it contains. The vector object is stored on the stack and contains a pointer to the content (`[1, 2, 3]`) stored on the heap. When we move `v` to `v2`, it creates a copy of that pointer, for `v2`. Which means that there would be two pointers to the content of the vector on the heap. It would violate Rust's safety guarantees by introducing a data race. Therefore, Rust forbids using `v` after we've done the move.

It's also important to note that optimizations may remove the actual copy of the bytes on the stack, depending on circumstances. So it may not be as inefficient as it initially seems.

Copy types

We've established that when ownership is transferred to another binding, you cannot use the original binding. However, there's a trait that changes this behavior, and it's called `Copy`. We haven't discussed traits yet, but for now, you can think of them as an annotation to a particular type that adds extra behavior. For example:

```
let v = 1;  
  
let v2 = v;  
  
println!("v is: {}", v);
```

In this case, `v` is an `i32`, which implements the `Copy` trait. This means that, just like a move, when we assign `v` to `v2`, a copy of the data is made. But, unlike a move, we can still use `v` afterward. This is because an `i32` has no pointers to data somewhere else, copying it is a full copy.

All primitive types implement the `Copy` trait and their ownership is therefore not moved like one would assume, following the 'ownership rules'. To give an example, the two following snippets of code only compile because the `i32` and `bool` types implement the `Copy` trait.

```
fn main() {
    let a = 5;

    let _y = double(a);
    println!("{}", a);
}

fn double(x: i32) -> i32 {
    x * 2
}

fn main() {
    let a = true;

    let _y = change_truth(a);
    println!("{}", a);
}

fn change_truth(x: bool) -> bool {
    !x
}
```

If we had used types that do not implement the `Copy` trait, we would have gotten a compile error because we tried to use a moved value.

```
error: use of moved value: 'a'
println!("{}", a);
           ^
```

We will discuss how to make your own types `Copy` in the traits section.

More than ownership

Of course, if we had to hand ownership back with every function we wrote:

```
fn foo(v: Vec<i32>) -> Vec<i32> {
    // do stuff with v

    // hand back ownership
    v
}
```

This would get very tedious. It gets worse the more things we want to take ownership of:

```
fn foo(v1: Vec<i32>, v2: Vec<i32>) -> (Vec<i32>, Vec<i32>, i32) {  
    // do stuff with v1 and v2  
  
    // hand back ownership, and the result of our function  
    (v1, v2, 42)  
}  
  
let v1 = vec![1, 2, 3];  
let v2 = vec![1, 2, 3];  
  
let (v1, v2, answer) = foo(v1, v2);
```

Ugh! The return type, return line, and calling the function gets way more complicated.

Luckily, Rust offers a feature, borrowing, which helps us solve this problem. It's the topic of the next section!

4.8 References and Borrowing

This guide is one of three presenting Rust's ownership system. This is one of Rust's most unique and compelling features, with which Rust developers should become quite acquainted. Ownership is how Rust achieves its largest goal, memory safety. There are a few distinct concepts, each with its own chapter:

- ownership, which you're reading now
- borrowing, and their associated feature 'references'
- lifetimes, an advanced concept of borrowing

These three chapters are related, and in order. You'll need all three to fully understand the ownership system.

Meta

Before we get to the details, two important notes about the ownership system.

Rust has a focus on safety and speed. It accomplishes these goals through many 'zero-cost abstractions', which means that in Rust, abstractions cost as little as possible in order to make them work. The ownership system is a prime example of a zero-cost abstraction. All of the analysis we'll talk about in this guide is *done at compile time*. You do not pay any run-time cost for any of these features.

However, this system does have a certain cost: learning curve. Many new users to Rust experience something we like to call 'fighting with the borrow checker', where the Rust compiler refuses to compile a program that the author thinks is valid. This often happens because the programmer's mental model of how ownership should work doesn't

match the actual rules that Rust implements. You probably will experience similar things at first. There is good news, however: more experienced Rust developers report that once they work with the rules of the ownership system for a period of time, they fight the borrow checker less and less.

With that in mind, let's learn about borrowing.

Borrowing

At the end of the [Ownership](#) section, we had a nasty function that looked like this:

```
fn foo(v1: Vec<i32>, v2: Vec<i32>) -> (Vec<i32>, Vec<i32>, i32) {  
    // do stuff with v1 and v2  
  
    // hand back ownership, and the result of our function  
    (v1, v2, 42)  
}  
  
let v1 = vec![1, 2, 3];  
let v2 = vec![1, 2, 3];  
  
let (v1, v2, answer) = foo(v1, v2);
```

This is not idiomatic Rust, however, as it doesn't take advantage of borrowing. Here's the first step:

```
fn foo(v1: &Vec<i32>, v2: &Vec<i32>) -> i32 {  
    // do stuff with v1 and v2  
  
    // return the answer  
    42  
}  
  
let v1 = vec![1, 2, 3];  
let v2 = vec![1, 2, 3];  
  
let answer = foo(&v1, &v2);  
  
// we can use v1 and v2 here!
```

Instead of taking `Vec<i32>`s as our arguments, we take a reference: `&Vec<i32>`. And instead of passing `v1` and `v2` directly, we pass `&v1` and `&v2`. We call the `&T` type a 'reference', and rather than owning the resource, it borrows ownership. A binding that borrows something does not deallocate the resource when it goes out of scope. This means that after the call to `foo()`, we can use our original bindings again.

References are immutable, like bindings. This means that inside of `foo()`, the vectors can't be changed at all:

```
fn foo(v: &Vec<i32>) {
    v.push(5);
}

let v = vec![];

foo(&v);
```

errors with:

```
error: cannot borrow immutable borrowed content `*v` as mutable
v.push(5);
^
```

Pushing a value mutates the vector, and so we aren't allowed to do it.

&mut references

There's a second kind of reference: `&mut T`. A 'mutable reference' allows you to mutate the resource you're borrowing. For example:

```
let mut x = 5;
{
    let y = &mut x;
    *y += 1;
}
println!("{}", x);
```

This will print `6`. We make `y` a mutable reference to `x` then add one to the thing `y` points at. You'll notice that `x` had to be marked `mut` as well. If it wasn't, we couldn't take a mutable borrow to an immutable value.

You'll also notice we added an asterisk (*) in front of `y`, making it `*y`, this is because `y` is a `&mut` reference. You'll also need to use them for accessing the contents of a reference as well.

Otherwise, `&mut` references are like references. There is a large difference between the two, and how they interact, though. You can tell something is fishy in the above example, because we need that extra scope, with the `{` and `}`. If we remove them, we get an error:

```
error: cannot borrow `x` as immutable because it is also borrowed as mutable
println!("{}", x);
^

note: previous borrow of `x` occurs here; the mutable borrow prevents
subsequent moves, borrows, or modification of `x` until the borrow ends
    let y = &mut x;
    ^

note: previous borrow ends here
fn main() {
```

```
}  
^
```

As it turns out, there are rules.

The Rules

Here's the rules about borrowing in Rust:

First, any borrow must last for a scope no greater than that of the owner. Second, you may have one or the other of these two kinds of borrows, but not both at the same time:

- one or more references (**&T**) to a resource,
- exactly one mutable reference (**&mut T**).

You may notice that this is very similar, though not exactly the same as, to the definition of a data race:

There is a 'data race' when two or more pointers access the same memory location at the same time, where at least one of them is writing, and the operations are not synchronized.

With references, you may have as many as you'd like, since none of them are writing. However, as we can only have one **&mut** at a time, it is impossible to have a data race. This is how Rust prevents data races at compile time: we'll get errors if we break the rules.

With this in mind, let's consider our example again.

Thinking in scopes

Here's the code:

```
let mut x = 5;  
let y = &mut x;  
  
*y += 1;  
  
println!("{}", x);
```

This code gives us this error:

```
error: cannot borrow 'x' as immutable because it is also borrowed as mutable  
    println!("{}", x);  
                ^
```

This is because we've violated the rules: we have a **&mut T** pointing to **x**, and so we aren't allowed to create any **&Ts**. One or the other. The note hints at how to think about this problem:

```

note: previous borrow ends here
fn main() {

}
~

```

In other words, the mutable borrow is held through the rest of our example. What we want is for the mutable borrow to end before we try to call `println!` and make an immutable borrow. In Rust, borrowing is tied to the scope that the borrow is valid for. And our scopes look like this:

```

let mut x = 5;

let y = &mut x;    // -+ &mut borrow of x starts here
                // |
*y += 1;           // |
                // |
println!("{}", x); // -+ - try to borrow x here
                // -+ &mut borrow of x ends here

```

The scopes conflict: we can't make an `&x` while `y` is in scope.

So when we add the curly braces:

```

let mut x = 5;

{
    let y = &mut x; // -+ &mut borrow starts here
    *y += 1;       // |
}                 // -+ ... and ends here

println!("{}", x); // <- try to borrow x here

```

There's no problem. Our mutable borrow goes out of scope before we create an immutable one. But scope is the key to seeing how long a borrow lasts for.

Issues borrowing prevents

Why have these restrictive rules? Well, as we noted, these rules prevent data races. What kinds of issues do data races cause? Here's a few.

Iterator invalidation

One example is 'iterator invalidation', which happens when you try to mutate a collection that you're iterating over. Rust's borrow checker prevents this from happening:


```
let mut v = vec![1, 2, 3];

for i in &v {
    println!("{}", i);
}
```

This prints out one through three. As we iterate through the vector, we're only given references to the elements. And `v` is itself borrowed as immutable, which means we can't change it while we're iterating:

```
let mut v = vec![1, 2, 3];

for i in &v {
    println!("{}", i);
    v.push(34);
}
```

Here's the error:

```
error: cannot borrow 'v' as mutable because it is also borrowed as immutable
      v.push(34);
      ^

note: previous borrow of 'v' occurs here; the immutable borrow prevents
subsequent moves or mutable borrows of 'v' until the borrow ends
for i in &v {
      ^

note: previous borrow ends here
for i in &v {
    println!("{}", i);
    v.push(34);
}
^
```

We can't modify `v` because it's borrowed by the loop.

use after free

References must not live longer than the resource they refer to. Rust will check the scopes of your references to ensure that this is true.

If Rust didn't check this property, we could accidentally use a reference which was invalid. For example:

```
let y: &i32;
{
    let x = 5;
    y = &x;
}

println!("{}", y);
```

We get this error:

```
error: 'x' does not live long enough
  y = &x;
    ^

note: reference must be valid for the block suffix following statement 0 at
2:16...
let y: &i32;
{
    let x = 5;
    y = &x;
}

note: ...but borrowed value is only valid for the block suffix following
statement 0 at 4:18
    let x = 5;
    y = &x;
}
```

In other words, **y** is only valid for the scope where **x** exists. As soon as **x** goes away, it becomes invalid to refer to it. As such, the error says that the borrow ‘doesn’t live long enough’ because it’s not valid for the right amount of time.

The same problem occurs when the reference is declared before the variable it refers to. This is because resources within the same scope are freed in the opposite order they were declared:

```
let y: &i32;
let x = 5;
y = &x;

println!("{}", y);
```

We get this error:

```
error: 'x' does not live long enough
y = &x;
    ^

note: reference must be valid for the block suffix following statement 0 at
2:16...
    let y: &i32;
    let x = 5;
    y = &x;

    println!("{}", y);
}

note: ...but borrowed value is only valid for the block suffix following
statement 1 at 3:14
    let x = 5;
    y = &x;
```

```
println!("{}", y);  
}
```

In the above example, `y` is declared before `x`, meaning that `y` lives longer than `x`, which is not allowed.

4.9 Lifetimes

This guide is one of three presenting Rust's ownership system. This is one of Rust's most unique and compelling features, with which Rust developers should become quite acquainted. Ownership is how Rust achieves its largest goal, memory safety. There are a few distinct concepts, each with its own chapter:

- ownership, which you're reading now
- borrowing, and their associated feature 'references'
- lifetimes, an advanced concept of borrowing

These three chapters are related, and in order. You'll need all three to fully understand the ownership system.

Meta

Before we get to the details, two important notes about the ownership system.

Rust has a focus on safety and speed. It accomplishes these goals through many 'zero-cost abstractions', which means that in Rust, abstractions cost as little as possible in order to make them work. The ownership system is a prime example of a zero-cost abstraction. All of the analysis we'll talk about in this guide is *done at compile time*. You do not pay any run-time cost for any of these features.

However, this system does have a certain cost: learning curve. Many new users to Rust experience something we like to call 'fighting with the borrow checker', where the Rust compiler refuses to compile a program that the author thinks is valid. This often happens because the programmer's mental model of how ownership should work doesn't match the actual rules that Rust implements. You probably will experience similar things at first. There is good news, however: more experienced Rust developers report that once they work with the rules of the ownership system for a period of time, they fight the borrow checker less and less.

With that in mind, let's learn about lifetimes.

Lifetimes

Lending out a reference to a resource that someone else owns can be complicated. For example, imagine this set of operations:

1. I acquire a handle to some kind of resource.
2. I lend you a reference to the resource.
3. I decide I'm done with the resource, and deallocate it, while you still have your reference.
4. You decide to use the resource.

Uh oh! Your reference is pointing to an invalid resource. This is called a dangling pointer or 'use after free', when the resource is memory.

To fix this, we have to make sure that step four never happens after step three. The ownership system in Rust does this through a concept called lifetimes, which describe the scope that a reference is valid for.

When we have a function that takes a reference by argument, we can be implicit or explicit about the lifetime of the reference:

```
// implicit
fn foo(x: &i32) {
}

// explicit
fn bar<'a>(x: &'a i32) {
}
```

The `'a` reads 'the lifetime a'. Technically, every reference has some lifetime associated with it, but the compiler lets you elide (i.e. omit, see "[Lifetime Elision](#)" below) them in common cases. Before we get to that, though, let's break the explicit example down:

```
fn bar<'a>(...)
```

We previously talked a little about function syntax ([Functions](#)), but we didn't discuss the `<>`s after a function's name. A function can have 'generic parameters' between the `<>`s, of which lifetimes are one kind. We'll discuss other kinds of generics later in the book, but for now, let's focus on the lifetimes aspect.

We use `<>` to declare our lifetimes. This says that `bar` has one lifetime, `'a`. If we had two reference parameters, it would look like this:

```
fn bar<'a, 'b>(...)
```

Then in our parameter list, we use the lifetimes we've named:

```
...(x: &'a i32)
```

If we wanted a `&mut` reference, we'd do this:

```
...(x: &'a mut i32)
```

If you compare `&mut i32` to `&'a mut i32`, they're the same, it's that the lifetime `'a` has snuck in between the `&` and the `mut i32`. We read `&mut i32` as 'a mutable reference to an `i32`' and `&'a mut i32` as 'a mutable reference to an `i32` with the lifetime `'a`'.

In structs

You'll also need explicit lifetimes when working with structs that contain references:

```
struct Foo<'a> {  
    x: &'a i32,  
}  
  
fn main() {  
    let y = &5; // this is the same as 'let _y = 5; let y = &_y;'  
    let f = Foo { x: y };  
  
    println!("{}", f.x);  
}
```

As you can see, `structs` can also have lifetimes. In a similar way to functions,

```
struct Foo<'a> {
```

declares a lifetime, and

```
x: &'a i32,
```

uses it. So why do we need a lifetime here? We need to ensure that any reference to a `Foo` cannot outlive the reference to an `i32` it contains.

impl blocks

Let's implement a method on `Foo`:

```

struct Foo<'a> {
    x: &'a i32,
}

impl<'a> Foo<'a> {
    fn x(&self) -> &'a i32 { self.x }
}

fn main() {
    let y = &5; // this is the same as 'let _y = 5; let y = &_y;'
    let f = Foo { x: y };

    println!("x is: {}", f.x());
}

```

As you can see, we need to declare a lifetime for **Foo** in the **impl** line. We repeat **'a** twice, like on functions: **impl<'a>** defines a lifetime **'a**, and **Foo<'a>** uses it.

Multiple lifetimes

If you have multiple references, you can use the same lifetime multiple times:

```
fn x_or_y<'a>(x: &'a str, y: &'a str) -> &'a str {
```

This says that **x** and **y** both are alive for the same scope, and that the return value is also alive for that scope. If you wanted **x** and **y** to have different lifetimes, you can use multiple lifetime parameters:

```
fn x_or_y<'a, 'b>(x: &'a str, y: &'b str) -> &'a str {
```

In this example, **x** and **y** have different valid scopes, but the return value has the same lifetime as **x**.

Thinking in scopes

A way to think about lifetimes is to visualize the scope that a reference is valid for. For example:

```

fn main() {
    let y = &5;      // -+ y goes into scope
                    // |
    // stuff        // |
                    // |
                    // -+ y goes out of scope
}

```

Adding in our **Foo**:

```

struct Foo<'a> {
    x: &'a i32,
}

fn main() {
    let y = &5;           // -+ y goes into scope
    let f = Foo { x: y }; // -+ f goes into scope
    // stuff              // |
                          // |
                          // -+ f and y go out of scope
}

```

Our `f` lives within the scope of `y`, so everything works. What if it didn't? This code won't work:

```

struct Foo<'a> {
    x: &'a i32,
}

fn main() {
    let x;                // -+ x goes into scope
                          // |
    {                     // |
        let y = &5;       // ----+ y goes into scope
        let f = Foo { x: y }; // ----+ f goes into scope
        x = &f.x;         // | | error here
    }                     // ----+ f and y go out of scope
                          // |
    println!("{}", x);    // |
                          // -+ x goes out of scope
}

```

Whew! As you can see here, the scopes of `f` and `y` are smaller than the scope of `x`. But when we do `x = &f.x`, we make `x` a reference to something that's about to go out of scope.

Named lifetimes are a way of giving these scopes a name. Giving something a name is the first step towards being able to talk about it.

'static

The lifetime named 'static' is a special lifetime. It signals that something has the lifetime of the entire program. Most Rust programmers first come across 'static when dealing with strings:

```
let x: &'static str = "Hello, world.";
```

String literals have the type `&'static str` because the reference is always alive: they are baked into the data segment of the final binary. Another example are globals:

```
static F00: i32 = 5;  
let x: &'static i32 = &F00;
```

This adds an `i32` to the data segment of the binary, and `x` is a reference to it.

Lifetime Elision

Rust supports powerful local type inference in function bodies, but it's forbidden in item signatures to allow reasoning about the types based on the item signature alone. However, for ergonomic reasons a very restricted secondary inference algorithm called “lifetime elision” applies in function signatures. It infers only based on the signature components themselves and not based on the body of the function, only infers lifetime parameters, and does this with only three easily memorizable and unambiguous rules. This makes lifetime elision a shorthand for writing an item signature, while not hiding away the actual types involved as full local inference would if applied to it.

When talking about lifetime elision, we use the term *input lifetime* and *output lifetime*. An *input lifetime* is a lifetime associated with a parameter of a function, and an *output lifetime* is a lifetime associated with the return value of a function. For example, this function has an input lifetime:

```
fn foo<'a>(bar: &'a str)
```

This one has an output lifetime:

```
fn foo<'a>() -> &'a str
```

This one has a lifetime in both positions:

```
fn foo<'a>(bar: &'a str) -> &'a str
```

Here are the three rules:

- Each elided lifetime in a function's arguments becomes a distinct lifetime parameter.
- If there is exactly one input lifetime, elided or not, that lifetime is assigned to all elided lifetimes in the return values of that function.
- If there are multiple input lifetimes, but one of them is `&self` or `&mut self`, the lifetime of self is assigned to all elided output lifetimes.

Otherwise, it is an error to elide an output lifetime.

Examples

Here are some examples of functions with elided lifetimes. We've paired each example of an elided lifetime with its expanded form.

```
fn print(s: &str); // elided
fn print<'a>(s: &'a str); // expanded

fn debug(lvl: u32, s: &str); // elided
fn debug<'a>(lvl: u32, s: &'a str); // expanded

// In the preceding example, 'lvl' doesn't need a lifetime because it's not a
// reference ('&'). Only things relating to references (such as a 'struct'
// which contains a reference) need lifetimes.

fn substr(s: &str, until: u32) -> &str; // elided
fn substr<'a>(s: &'a str, until: u32) -> &'a str; // expanded

fn get_str() -> &str; // ILLEGAL, no inputs

fn frob(s: &str, t: &str) -> &str; // ILLEGAL, two inputs
fn frob<'a, 'b>(s: &'a str, t: &'b str) -> &str; // Expanded: Output lifetime is
    ↪ ambiguous

fn get_mut(&mut self) -> &mut T; // elided
fn get_mut<'a>(&'a mut self) -> &'a mut T; // expanded

fn args<T: ToCStr>(&mut self, args: &[T]) -> &mut Command; // elided
fn args<'a, 'b, T: ToCStr>(&'a mut self, args: &'b [T]) -> &'a mut Command; //
    ↪ expanded

fn new(buf: &mut [u8]) -> BufWriter; // elided
fn new<'a>(buf: &'a mut [u8]) -> BufWriter<'a>; // expanded
```

4.10 Mutability

Mutability, the ability to change something, works a bit differently in Rust than in other languages. The first aspect of mutability is its non-default status:

```
let x = 5;
x = 6; // error!
```

We can introduce mutability with the `mut` keyword:

```
let mut x = 5;

x = 6; // no problem!
```

This is a mutable [Variable Bindings](#). When a binding is mutable, it means you're allowed to change what the binding points to. So in the above example, it's not so much

that the value at `x` is changing, but that the binding changed from one `i32` to another.

If you want to change what the binding points to, you'll need a mutable reference (see [References and Borrowing](#)):

```
let mut x = 5;
let y = &mut x;
```

`y` is an immutable binding to a mutable reference, which means that you can't bind `y` to something else (`y = &mut z`), but you can mutate the thing that's bound to `y` (`*y = 5`). A subtle distinction.

Of course, if you need both:

```
let mut x = 5;
let mut y = &mut x;
```

Now `y` can be bound to another value, and the value it's referencing can be changed.

It's important to note that `mut` is part of a pattern, so you can do things like this:

```
let (mut x, y) = (5, 6);

fn foo(mut x: i32) {
```

Interior vs. Exterior Mutability

However, when we say something is 'immutable' in Rust, that doesn't mean that it's not able to be changed: we mean something has 'exterior mutability'. Consider, for example, `Arc<T>`:

```
use std::sync::Arc;

let x = Arc::new(5);
let y = x.clone();
```

When we call `clone()`, the `Arc<T>` needs to update the reference count. Yet we've not used any muts here, `x` is an immutable binding, and we didn't take `&mut 5` or anything. So what gives?

To understand this, we have to go back to the core of Rust's guiding philosophy, memory safety, and the mechanism by which Rust guarantees it, the ownership system (see [Ownership](#)), and more specifically, borrowing (see [References and Borrowing](#)):

You may have one or the other of these two kinds of borrows, but not both at the same time:

- one or more references (`&T`) to a resource,
- exactly one mutable reference (`&mut T`).

So, that's the real definition of 'immutability': is this safe to have two pointers to? In `Arc<T>`'s case, yes: the mutation is entirely contained inside the structure itself. It's not user facing. For this reason, it hands out `&T` with `clone()`. If it handed out `&mut T`s, though, that would be a problem.

Other types, like the ones in the `std::cell` module, have the opposite: interior mutability. For example:

```
use std::cell::RefCell;

let x = RefCell::new(42);

let y = x.borrow_mut();
```

`RefCell` hands out `&mut` references to what's inside of it with the `borrow_mut()` method. Isn't that dangerous? What if we do:

```
use std::cell::RefCell;

let x = RefCell::new(42);

let y = x.borrow_mut();
let z = x.borrow_mut();
```

This will in fact panic, at runtime. This is what `RefCell` does: it enforces Rust's borrowing rules at runtime, and `panic!`s if they're violated. This allows us to get around another aspect of Rust's mutability rules. Let's talk about it first.

Field-level mutability

Mutability is a property of either a borrow (`&mut`) or a binding (`let mut`). This means that, for example, you cannot have a struct with some fields mutable and some immutable:

```
struct Point {
    x: i32,
    mut y: i32, // nope
}
```

The mutability of a struct is in its binding:

```

struct Point {
    x: i32,
    y: i32,
}

let mut a = Point { x: 5, y: 6 };

a.x = 10;

let b = Point { x: 5, y: 6 };

b.x = 10; // error: cannot assign to immutable field 'b.x'

```

However, by using `Cell<T>`, you can emulate field-level mutability:

```

use std::cell::Cell;

struct Point {
    x: i32,
    y: Cell<i32>,
}

let point = Point { x: 5, y: Cell::new(6) };

point.y.set(7);

println!("y: {:?}", point.y);

```

This will print `y: Cell value: 7`. We've successfully updated `y`.

4.11 Structs

structs are a way of creating more complex data types. For example, if we were doing calculations involving coordinates in 2D space, we would need both an `x` and a `y` value:

```

let origin_x = 0;
let origin_y = 0;

```

A **struct** lets us combine these two into a single, unified datatype with `x` and `y` as field labels:

```

struct Point {
    x: i32,
    y: i32,
}

fn main() {
    let origin = Point { x: 0, y: 0 }; // origin: Point

    println!("The origin is at ({}, {})", origin.x, origin.y);
}

```

There's a lot going on here, so let's break it down. We declare a `struct` with the `struct` keyword, and then with a name. By convention, `structs` begin with a capital letter and are camel cased: `PointInSpace`, not `Point_In_Space`.

We can create an instance of our `struct` via `let`, as usual, but we use a `key: value` style syntax to set each field. The order doesn't need to be the same as in the original declaration.

Finally, because fields have names, we can access them through dot notation: `origin.x`.

The values in `structs` are immutable by default, like other bindings in Rust. Use `mut` to make them mutable:

```
struct Point {
    x: i32,
    y: i32,
}

fn main() {
    let mut point = Point { x: 0, y: 0 };

    point.x = 5;

    println!("The point is at ({}, {})", point.x, point.y);
}
```

This will print `The point is at (5, 0)`.

Rust does not support field mutability at the language level, so you cannot write something like this:

```
struct Point {
    mut x: i32,
    y: i32,
}
```

Mutability is a property of the binding, not of the structure itself. If you're used to field-level mutability, this may seem strange at first, but it significantly simplifies things. It even lets you make things mutable on a temporary basis:

```
struct Point {
    x: i32,
    y: i32,
}

fn main() {
    let mut point = Point { x: 0, y: 0 };

    point.x = 5;

    let point = point; // now immutable

    point.y = 6; // this causes an error
}
```

Your structure can still contain `&mut` pointers, which will let you do some kinds of mutation:

```
struct Point {
    x: i32,
    y: i32,
}

struct PointRef<'a> {
    x: &'a mut i32,
    y: &'a mut i32,
}

fn main() {
    let mut point = Point { x: 0, y: 0 };

    {
        let r = PointRef { x: &mut point.x, y: &mut point.y };

        *r.x = 5;
        *r.y = 6;
    }

    assert_eq!(5, point.x);
    assert_eq!(6, point.y);
}
```

Update syntax

A `struct` can include `..` to indicate that you want to use a copy of some other `struct` for some of the values. For example:

```
struct Point3d {  
    x: i32,  
    y: i32,  
    z: i32,  
}  
  
let mut point = Point3d { x: 0, y: 0, z: 0 };  
point = Point3d { y: 1, .. point };
```

This gives `point` a new `y`, but keeps the old `x` and `z` values. It doesn't have to be the same `struct` either, you can use this syntax when making new ones, and it will copy the values you don't specify:

```
let origin = Point3d { x: 0, y: 0, z: 0 };  
let point = Point3d { z: 1, x: 2, .. origin };
```

Tuple structs

Rust has another data type that's like a hybrid between a tuple and a `struct`, called a 'tuple struct'. Tuple structs have a name, but their fields don't. They are declared with the `struct` keyword, and then with a name followed by a tuple:

```
struct Color(i32, i32, i32);  
struct Point(i32, i32, i32);  
  
let black = Color(0, 0, 0);  
let origin = Point(0, 0, 0);
```

Here, `black` and `origin` are not equal, even though they contain the same values.

It is almost always better to use a `struct` than a tuple struct. We would write `Color` and `Point` like this instead:

```
struct Color {  
    red: i32,  
    blue: i32,  
    green: i32,  
}  
  
struct Point {  
    x: i32,  
    y: i32,  
    z: i32,  
}
```

Good names are important, and while values in a tuple struct can be referenced with dot notation as well, a `struct` gives us actual names, rather than positions.

There is one case when a tuple struct is very useful, though, and that is when it has

only one element. We call this the ‘newtype’ pattern, because it allows you to create a new type that is distinct from its contained value and also expresses its own semantic meaning:

```
struct Inches(i32);

let length = Inches(10);

let Inches(integer_length) = length;
println!("length is {} inches", integer_length);
```

As you can see here, you can extract the inner integer type through a destructuring `let`, as with regular tuples. In this case, the `let Inches(integer_length)` assigns `10` to `integer_length`.

Unit-like structs

You can define a `struct` with no members at all:

```
struct Electron;

let x = Electron;
```

Such a `struct` is called ‘unit-like’ because it resembles the empty tuple, `()`, sometimes called ‘unit’. Like a tuple struct, it defines a new type.

This is rarely useful on its own (although sometimes it can serve as a marker type), but in combination with other features, it can become useful. For instance, a library may ask you to create a structure that implements a certain trait to handle events. If you don’t have any data you need to store in the structure, you can create a unit-like `struct`.

4.12 Enums

An `enum` in Rust is a type that represents data that is one of several possible variants. Each variant in the `enum` can optionally have data associated with it:

```
enum Message {
    Quit,
    ChangeColor(i32, i32, i32),
    Move { x: i32, y: i32 },
    Write(String),
}
```

The syntax for defining variants resembles the syntaxes used to define structs: you can have variants with no data (like unit-like structs), variants with named data, and variants with unnamed data (like tuple structs). Unlike separate struct definitions, however, an `enum` is a single type. A value of the enum can match any of the variants.

For this reason, an enum is sometimes called a ‘sum type’: the set of possible values of the enum is the sum of the sets of possible values for each variant.

We use the `::` syntax to use the name of each variant: they’re scoped by the name of the `enum` itself. This allows both of these to work:

```
let x: Message = Message::Move { x: 3, y: 4 };

enum BoardGameTurn {
    Move { squares: i32 },
    Pass,
}

let y: BoardGameTurn = BoardGameTurn::Move { squares: 1 };
```

Both variants are named `Move`, but since they’re scoped to the name of the `enum` they can both be used without conflict.

A value of an `enum` type contains information about which variant it is, in addition to any data associated with that variant. This is sometimes referred to as a ‘tagged union’, since the data includes a ‘tag’ indicating what type it is. The compiler uses this information to enforce that you’re accessing the data in the enum safely. For instance, you can’t simply try to destructure a value as if it were one of the possible variants:

```
fn process_color_change(msg: Message) {
    let Message::ChangeColor(r, g, b) = msg; // compile-time error
}
```

Not supporting these operations may seem rather limiting, but it’s a limitation which we can overcome. There are two ways: by implementing equality ourselves, or by pattern matching variants with `match` expressions, which you’ll learn in the next section. We don’t know enough about Rust to implement equality yet, but we’ll find out in the traits section.

Constructors as functions

An `enum` constructor can also be used like a function. For example:

```
let m = Message::Write("Hello, world".to_string());
```

is the same as

```
fn foo(x: String) -> Message {
    Message::Write(x)
}

let x = foo("Hello, world".to_string());
```

This is not immediately useful to us, but when we get to closures, we’ll talk about

passing functions as arguments to other functions. For example, with iterators, we can do this to convert a vector of `Strings` into a vector of `Message::Writes`:

```
let v = vec!["Hello".to_string(), "World".to_string()];

let v1: Vec<Message> = v.into_iter().map(Message::Write).collect();
```

4.13 Match

Often, a simple `if/else` (see [If](#)) isn't enough, because you have more than two possible options. Also, conditions can get quite complex. Rust has a keyword, `match`, that allows you to replace complicated `if/else` groupings with something more powerful. Check it out:

```
let x = 5;

match x {
    1 => println!("one"),
    2 => println!("two"),
    3 => println!("three"),
    4 => println!("four"),
    5 => println!("five"),
    _ => println!("something else"),
}
```

`match` takes an expression and then branches based on its value. Each 'arm' of the branch is of the form `val => expression`. When the value matches, that arm's expression will be evaluated. It's called `match` because of the term 'pattern matching', which `match` is an implementation of. There's a separate section on patterns that covers all the patterns that are possible here.

One of the many advantages of `match` is it enforces 'exhaustiveness checking'. For example if we remove the last arm with the underscore `_`, the compiler will give us an error:

```
error: non-exhaustive patterns: '_' not covered
```

Rust is telling us that we forgot a value. The compiler infers from `x` that it can have any positive 32bit value; for example 1 to 2,147,483,647. The `_` acts as a 'catch-all', and will catch all possible values that aren't specified in an arm of `match`. As you can see with the previous example, we provide `match` arms for integers 1-5, if `x` is 6 or any other value, then it is caught by `_`.

`match` is also an expression, which means we can use it on the right-hand side of a `let` binding or directly where an expression is used:

```
let x = 5;

let number = match x {
    1 => "one",
    2 => "two",
    3 => "three",
    4 => "four",
    5 => "five",
    _ => "something else",
};
```

Sometimes it's a nice way of converting something from one type to another; in this example the integers are converted to `String`.

Matching on enums

Another important use of the `match` keyword is to process the possible variants of an `enum`:

```
enum Message {
    Quit,
    ChangeColor(i32, i32, i32),
    Move { x: i32, y: i32 },
    Write(String),
}

fn quit() { /* ... */ }
fn change_color(r: i32, g: i32, b: i32) { /* ... */ }
fn move_cursor(x: i32, y: i32) { /* ... */ }

fn process_message(msg: Message) {
    match msg {
        Message::Quit => quit(),
        Message::ChangeColor(r, g, b) => change_color(r, g, b),
        Message::Move { x: x, y: y } => move_cursor(x, y),
        Message::Write(s) => println!("{}", s),
    };
}
```

Again, the Rust compiler checks exhaustiveness, so it demands that you have a match arm for every variant of the enum. If you leave one off, it will give you a compile-time error unless you use `_` or provide all possible arms.

Unlike the previous uses of `match`, you can't use the normal `if` statement to do this. You can use the `if let` (see) statement, which can be seen as an abbreviated form of `match`.

4.14 Patterns

Patterns are quite common in Rust. We use them in [Variable Bindings](#), `match` statements (see [Match](#)), and other places, too. Let's go on a whirlwind tour of all of the things patterns can do!

A quick refresher: you can match against literals directly, and `_` acts as an 'any' case:

```
let x = 1;

match x {
    1 => println!("one"),
    2 => println!("two"),
    3 => println!("three"),
    _ => println!("anything"),
}
```

This prints `one`.

There's one pitfall with patterns: like anything that introduces a new binding, they introduce shadowing. For example:

```
let x = 1;
let c = 'c';

match c {
    x => println!("x: {} c: {}", x, c),
}

println!("x: {}", x)
```

This prints:

```
x: c c: c
x: 1
```

In other words, `x =>` matches the pattern and introduces a new binding named `x`. This new binding is in scope for the match arm and takes on the value of `c`. Notice that the value of `x` outside the scope of the match has no bearing on the value of `x` within it. Because we already have a binding named `x` this new `x` shadows it.

Multiple patterns

You can match multiple patterns with `|`:

```
let x = 1;

match x {
    1 | 2 => println!("one or two"),
    3 => println!("three"),
    _ => println!("anything"),
}
```

This prints **one or two**.

Destructuring

If you have a compound data type, like a **struct** (see [Structs](#)), you can destructure it inside of a pattern:

```
struct Point {  
  x: i32,  
  y: i32,  
}  
  
let origin = Point { x: 0, y: 0 };  
  
match origin {  
  Point { x, y } => println!("{}", x, y),  
}
```

We can use **:** to give a value a different name.

```
struct Point {  
  x: i32,  
  y: i32,  
}  
  
let origin = Point { x: 0, y: 0 };  
  
match origin {  
  Point { x: x1, y: y1 } => println!("{}", x1, y1),  
}
```

If we only care about some of the values, we don't have to give them all names:

```
struct Point {  
  x: i32,  
  y: i32,  
}  
  
let origin = Point { x: 0, y: 0 };  
  
match origin {  
  Point { x, .. } => println!("x is {}", x),  
}
```

This prints **x is 0**.

You can do this kind of match on any member, not only the first:

```
struct Point {
    x: i32,
    y: i32,
}

let origin = Point { x: 0, y: 0 };

match origin {
    Point { y, .. } => println!("y is {}", y),
}
```

This prints `y is 0`.

This ‘destructuring’ behavior works on any compound data type, like tuples or enums.

Ignoring bindings

You can use `_` in a pattern to disregard the type and value. For example, here’s a `match` against a `Result<T, E>`:

```
match some_value {
    Ok(value) => println!("got a value: {}", value),
    Err(_) => println!("an error occurred"),
}
```

In the first arm, we bind the value inside the `Ok` variant to `value`. But in the `Err` arm, we use `_` to disregard the specific error, and print a general error message.

`_` is valid in any pattern that creates a binding. This can be useful to ignore parts of a larger structure:

```
fn coordinate() -> (i32, i32, i32) {
    // generate and return some sort of triple tuple
}

let (x, _, z) = coordinate();
```

Here, we bind the first and last element of the tuple to `x` and `z`, but ignore the middle element.

Similarly, you can use `..` in a pattern to disregard multiple values.

```
enum OptionalTuple {  
    Value(i32, i32, i32),  
    Missing,  
}  
  
let x = OptionalTuple::Value(5, -2, 3);  
  
match x {  
    OptionalTuple::Value(..) => println!("Got a tuple!"),  
    OptionalTuple::Missing => println!("No such luck."),  
}
```

This prints `Got a tuple!`.

ref and ref mut

If you want to get a reference (see [References and Borrowing](#)), use the `ref` keyword:

```
let x = 5;  
  
match x {  
    ref r => println!("Got a reference to {}", r),  
}
```

This prints `Got a reference to 5`.

Here, the `r` inside the `match` has the type `&i32`. In other words, the `ref` keyword creates a reference, for use in the pattern. If you need a mutable reference, `refmut` will work in the same way:

```
let mut x = 5;  
  
match x {  
    ref mut mr => println!("Got a mutable reference to {}", mr),  
}
```

Ranges

You can match a range of values with `...`:

```
let x = 1;  
  
match x {  
    1 ... 5 => println!("one through five"),  
    _ => println!("anything"),  
}
```

This prints `one through five`.

Ranges are mostly used with integers and `chars`:

```
let x = 'ä';

match x {
  'a' ... 'j' => println!("early letter"),
  'k' ... 'z' => println!("late letter"),
  _ => println!("something else"),
}
```

This prints `something else`.

Bindings

You can bind values to names with `@`:

```
let x = 1;

match x {
  e @ 1 ... 5 => println!("got a range element {}", e),
  _ => println!("anything"),
}
```

This prints `got a range element 1`. This is useful when you want to do a complicated match of part of a data structure:

```
#[derive(Debug)]
struct Person {
  name: Option<String>,
}

let name = "Steve".to_string();
let mut x: Option<Person> = Some(Person { name: Some(name) });
match x {
  Some(Person { name: ref a @ Some(_), .. }) => println!("{:?}", a),
  _ => {}
}
```

This prints `Some("Steve")`: we've bound the inner `name` to `a`.

If you use `@` with `|`, you need to make sure the name is bound in each part of the pattern:

```
let x = 5;

match x {
  e @ 1 ... 5 | e @ 8 ... 10 => println!("got a range element {}", e),
  _ => println!("anything"),
}
```


Guards

You can introduce ‘match guards’ with `if`:

```
enum OptionalInt {
    Value(i32),
    Missing,
}

let x = OptionalInt::Value(5);

match x {
    OptionalInt::Value(i) if i > 5 => println!("Got an int bigger than five!"),
    OptionalInt::Value(..) => println!("Got an int!"),
    OptionalInt::Missing => println!("No such luck."),
}
```

This prints `Got an int!`.

If you’re using `if` with multiple patterns, the `if` applies to both sides:

```
let x = 4;
let y = false;

match x {
    4 | 5 if y => println!("yes"),
    _ => println!("no"),
}
```

This prints `no`, because the `if` applies to the whole of `4 | 5`, and not to only the `5`. In other words, the precedence of `if` behaves like this:

```
(4 | 5) if y => ...
```

not this:

```
4 | (5 if y) => ...
```

Mix and Match

Whew! That’s a lot of different ways to match things, and they can all be mixed and matched, depending on what you’re doing:

```
match x {
    Foo { x: Some(ref name), y: None } => ...
}
```

Patterns are very powerful. Make good use of them.

5 Effective Rust

6 Nightly Rust

7 Glossary

Not every Rustacean has a background in systems programming, nor in computer science, so we've added explanations of terms that might be unfamiliar.

Abstract Syntax Tree

When a compiler is compiling your program, it does a number of different things. One of the things that it does is turn the text of your program into an 'abstract syntax tree', or 'AST'. This tree is a representation of the structure of your program. For example, `2 + 3` can be turned into a tree:



And `2 + (3 * 4)` would look like this:



Arity

Arity refers to the number of arguments a function or operation takes.

```
let x = (2, 3);
let y = (4, 6);
let z = (8, 2, 6);
```

In the example above `x` and `y` have arity 2. `z` has arity 3.

Bounds

Bounds are constraints on a type or trait. For example, if a bound is placed on the argument a function takes, types passed to that function must abide by that constraint.

DST (Dynamically Sized Type)

A type without a statically known size or alignment. ([more info](#))

Expression

In computer programming, an expression is a combination of values, constants, variables, operators and functions that evaluate to a single value. For example, `2 + (3 * 4)` is an expression that returns the value 14. It is worth noting that expressions can have

side-effects. For example, a function included in an expression might perform actions other than simply returning a value.

Expression-Oriented Language

In early programming languages, [Expression](#) and [Statement](#) were two separate syntactic categories: expressions had a value and statements did things. However, later languages blurred this distinction, allowing expressions to do things and statements to have a value. In an expression-oriented language, (nearly) every statement is an expression and therefore returns a value. Consequently, these expression statements can themselves form part of larger expressions.

Statement

In computer programming, a statement is the smallest standalone element of a programming language that commands a computer to perform an action.

8 Syntax Index

9 Bibliography

This is a reading list of material relevant to Rust. It includes prior research that has - at one time or another - influenced the design of Rust, as well as publications about Rust.

9.1 Type system

- [Region based memory management in Cyclone](#)
- [Safe manual memory management in Cyclone](#)
- [Typeclasses: making ad-hoc polymorphism less ad hoc](#)
- [Macros that work together](#)
- [Traits: composable units of behavior](#)
- [Alias burying](#) - We tried something similar and abandoned it.
- [External uniqueness is unique enough](#)
- [Uniqueness and Reference Immutability for Safe Parallelism](#)
- [Region Based Memory Management](#)

9.2 Concurrency

- [Singularity: rethinking the software stack](#)
- [Language support for fast and reliable message passing in singularity OS](#)
- [Scheduling multithreaded computations by work stealing](#)
- [Thread scheduling for multiprogramming multiprocessors](#)
- [The data locality of work stealing](#)
- [Dynamic circular work stealing deque](#) - The Chase/Lev deque
- [Work-first and help-first scheduling policies for async-finish task parallelism](#) - More general than fully-strict work stealing
- [A Java fork/join calamity](#) - critique of Java's fork/join library, particularly its application of work stealing to non-strict computation
- [Scheduling techniques for concurrent systems](#)
- [Contention aware scheduling](#)
- [Balanced work stealing for time-sharing multicores](#)

- [Three layer cake for shared-memory programming](#)
- [Non-blocking steal-half work queues](#)
- [Reagents: expressing and composing fine-grained concurrency](#)
- [Algorithms for scalable synchronization of shared-memory multiprocessors](#)
- [Epoch-based reclamation.](#)

9.3 Others

- [Crash-only software](#)
- [Composing High-Performance Memory Allocators](#)
- [Reconsidering Custom Memory Allocation](#)

9.4 Papers about Rust

- [GPU Programming in Rust: Implementing High Level Abstractions in a Systems Level Language.](#) Early GPU work by Eric Holk.
- [Parallel closures: a new twist on an old idea](#)
 - not exactly about Rust, but by nmatsakis
- [Patina: A Formalization of the Rust Programming Language.](#) Early formalization of a subset of the type system, by Eric Reed.
- [Experience Report: Developing the Servo Web Browser Engine using Rust.](#) By Lars Bergstrom.
- [Implementing a Generic Radix Trie in Rust.](#) Undergrad paper by Michael Sproul.
- [Reenix: Implementing a Unix-Like Operating System in Rust.](#) Undergrad paper by Alex Light.
- [Evaluation of performance and productivity metrics of potential programming languages in the HPC environment.](#) Bachelor's thesis by Florian Wilkens. Compares C, Go and Rust.
- [Nom, a byte oriented, streaming, zero copy, parser combinators library in Rust.](#) By Geoffroy Couprie, research for VLC.
- [Graph-Based Higher-Order Intermediate Representation.](#) An experimental IR implemented in Impala, a Rust-like language.
- [Code Refinement of Stencil Codes.](#) Another paper using Impala.

- [Parallelization in Rust with fork-join and friends](#). Linus Farnstrand's master's thesis.
- [Session Types for Rust](#). Philip Munksgaard's master's thesis. Research for Servo.
- [Ownership is Theft: Experiences Building an Embedded OS in Rust](#) - Amit Levy, et. al.